

How Google Does Machine Learning

What it Means to be AI-first

How Google Does ML

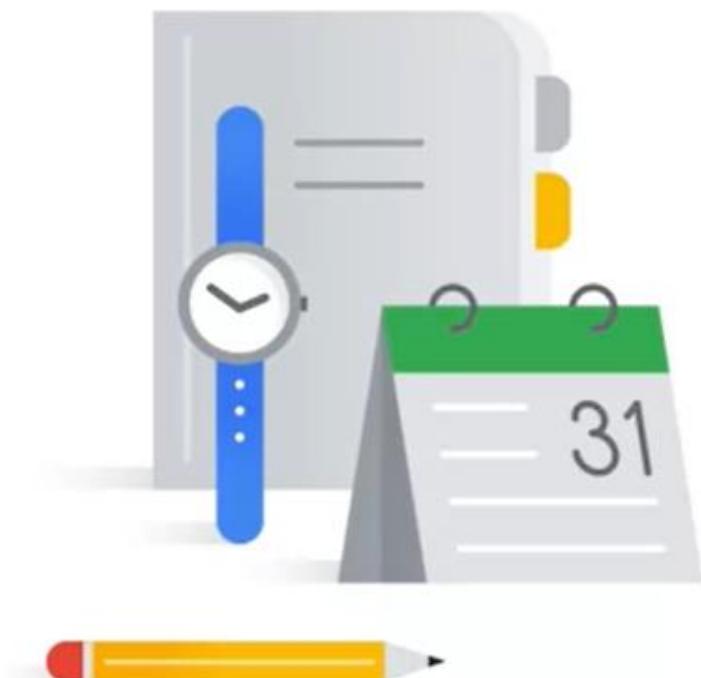
Introduction to Machine Learning Development with Vertex AI

ML Development with Vertex Notebooks

Best Practices for Implementing Machine Learning on Vertex AI

Responsible AI Development

Course Summary



Practical, real-world introduction to ML



Data Analysts
Citizen Data Scientists

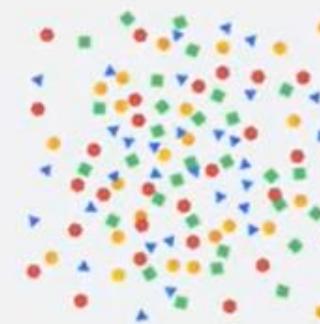
Build without writing a
single piece of code



ML Engineers
ML Scientists

Go “code deep” with
custom training

Machine learning is a way to use standard algorithms to derive predictive insights from data and make repeated decisions.



Data



Algorithm

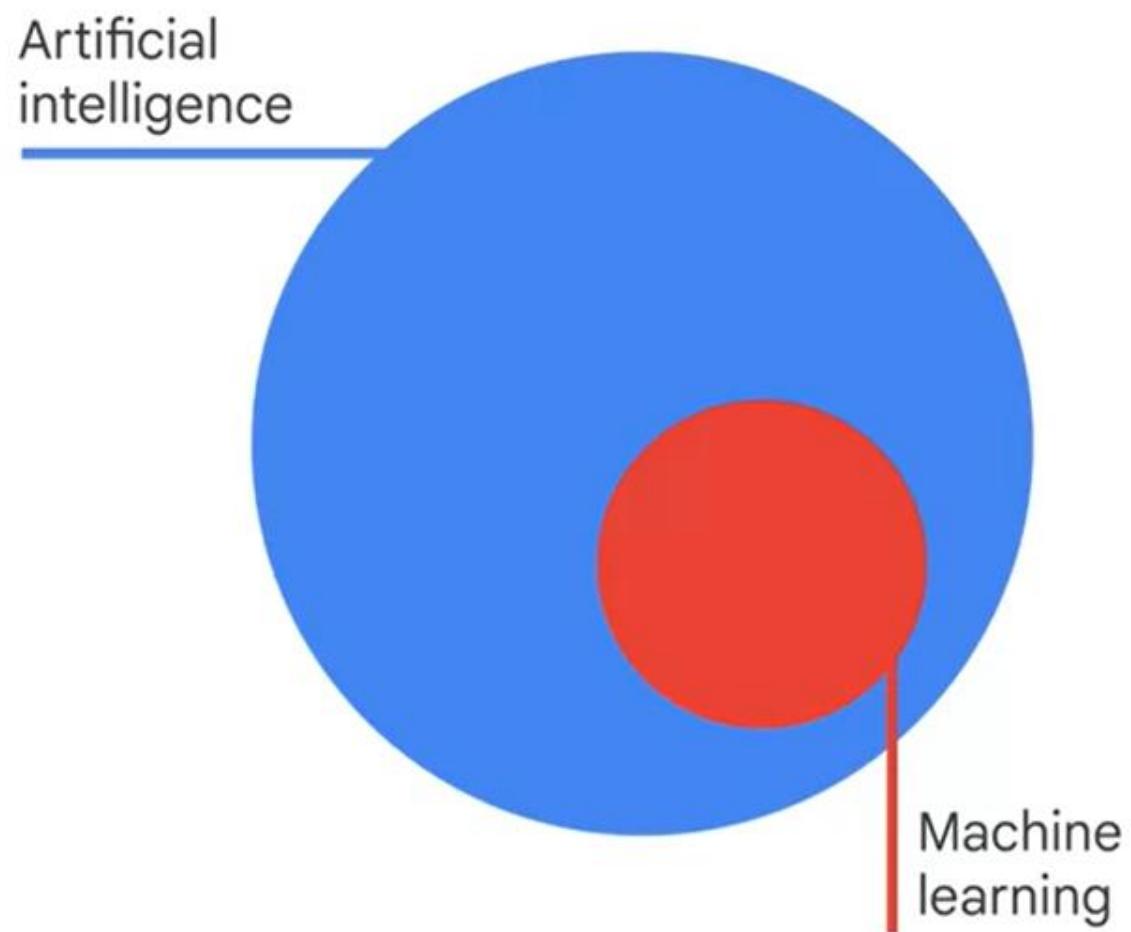


Predictive insight

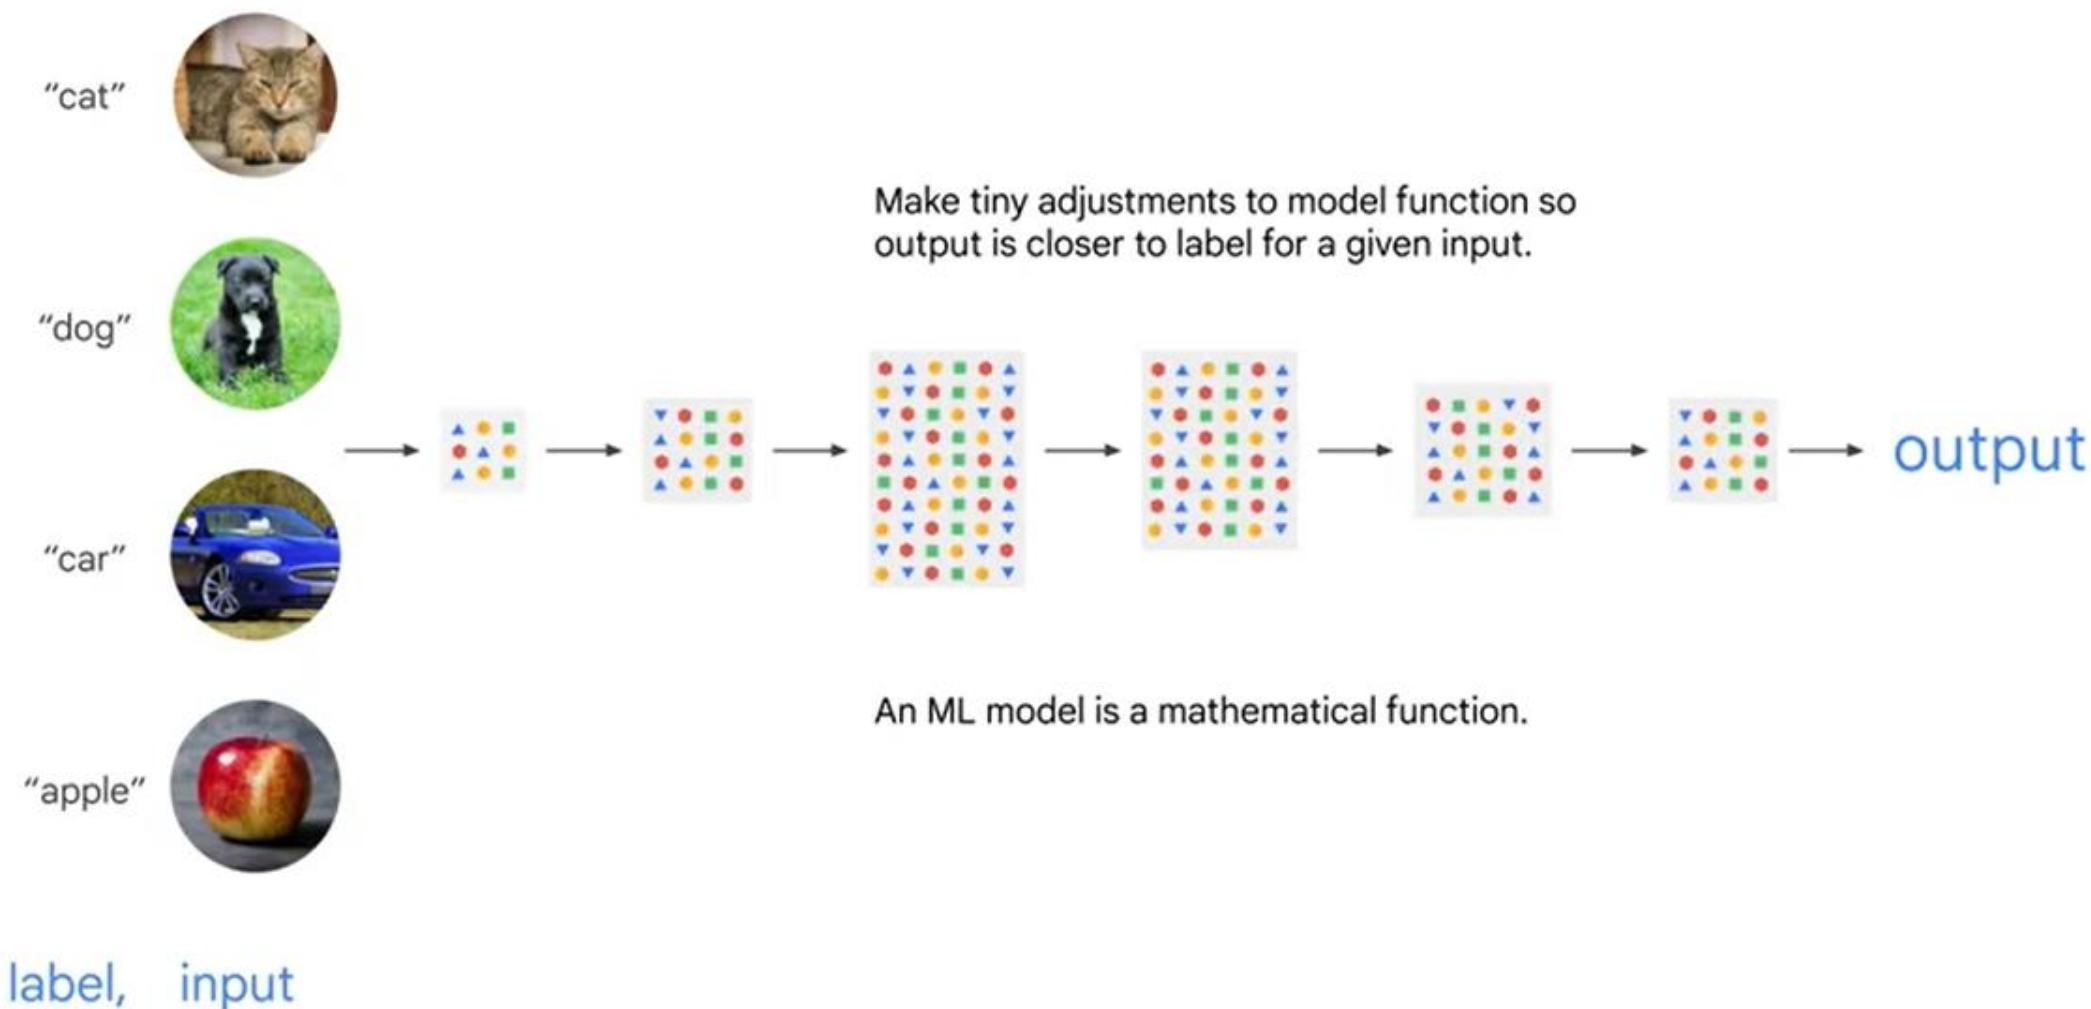


Decision

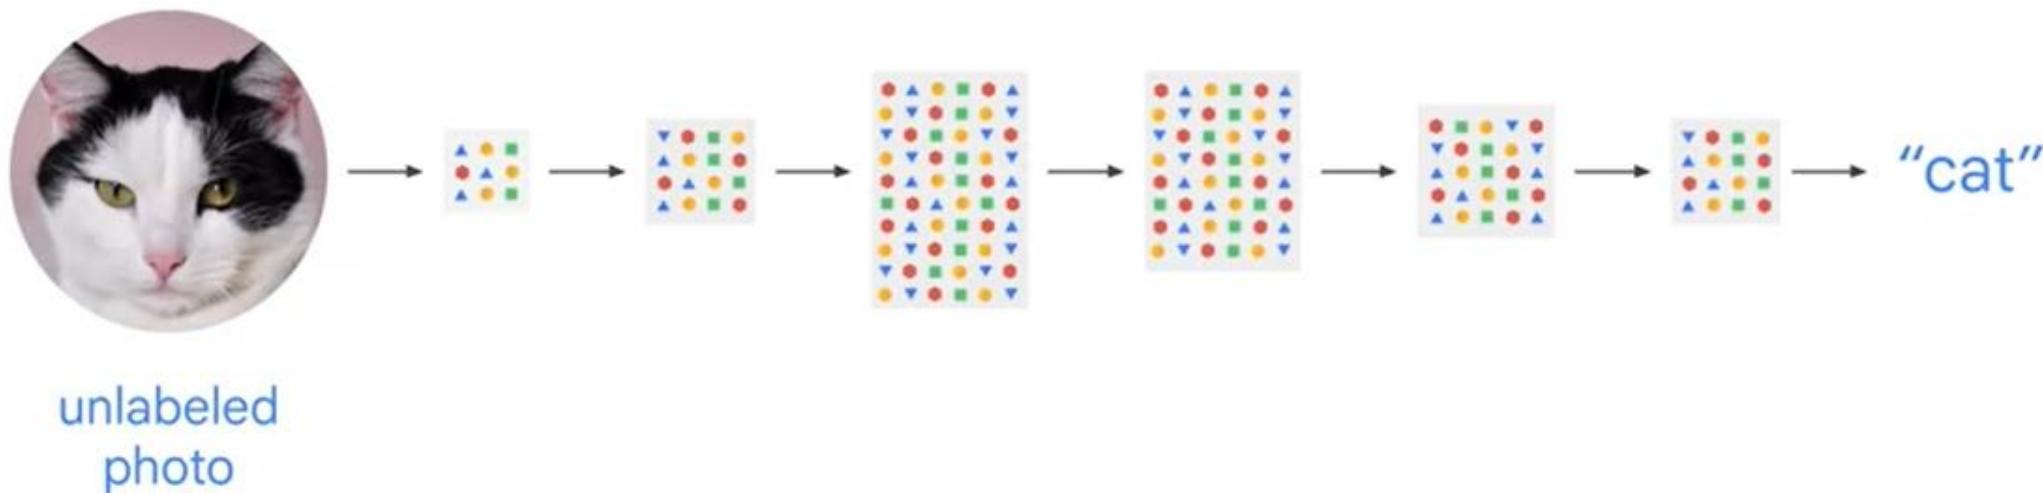
Artificial Intelligence is a discipline; machine learning is a specific way of solving AI problems.

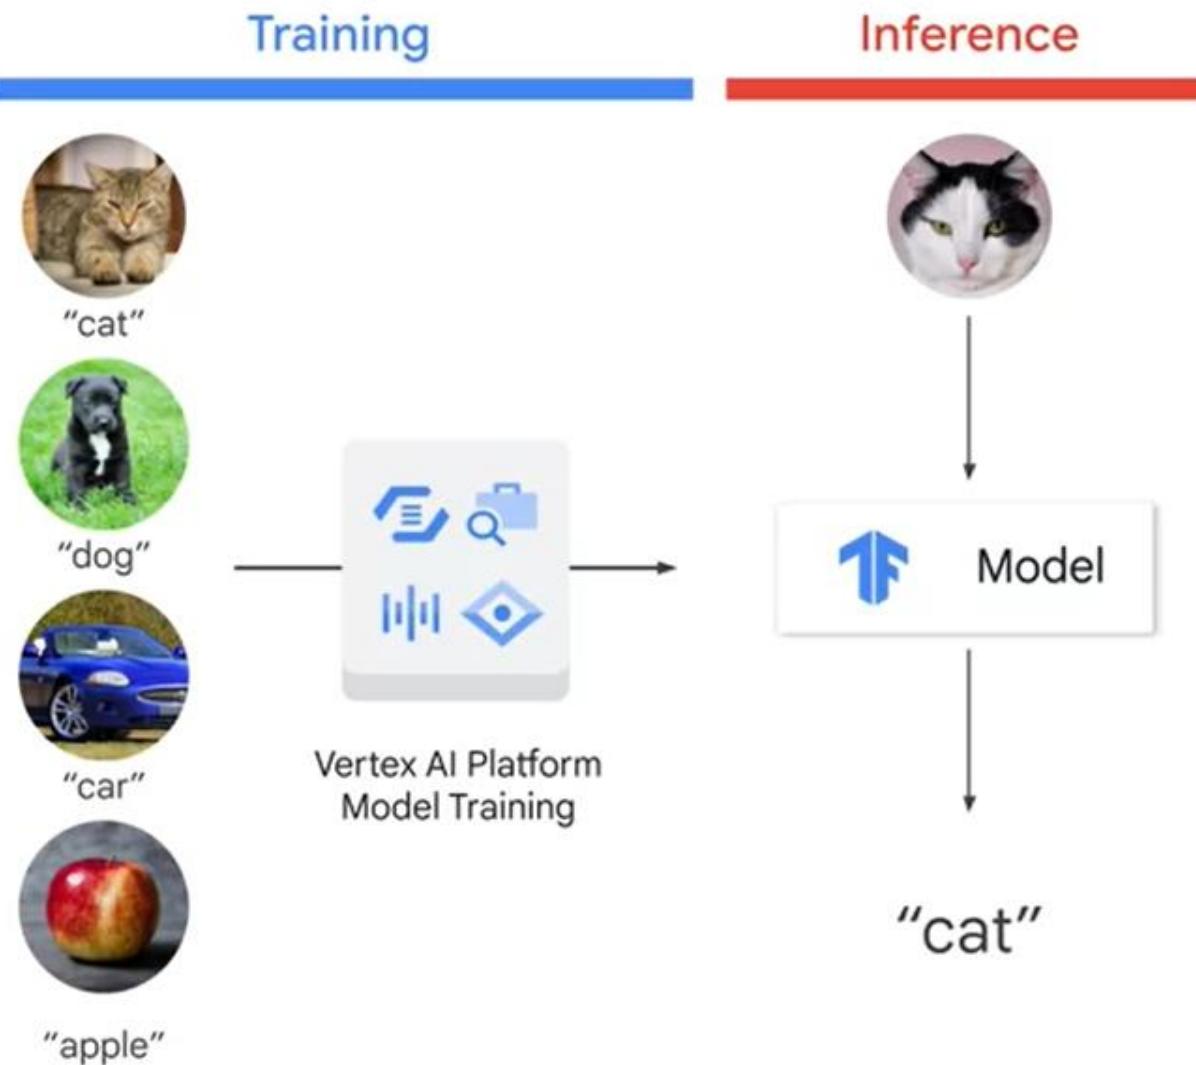


Stage 1: Train an ML model with examples



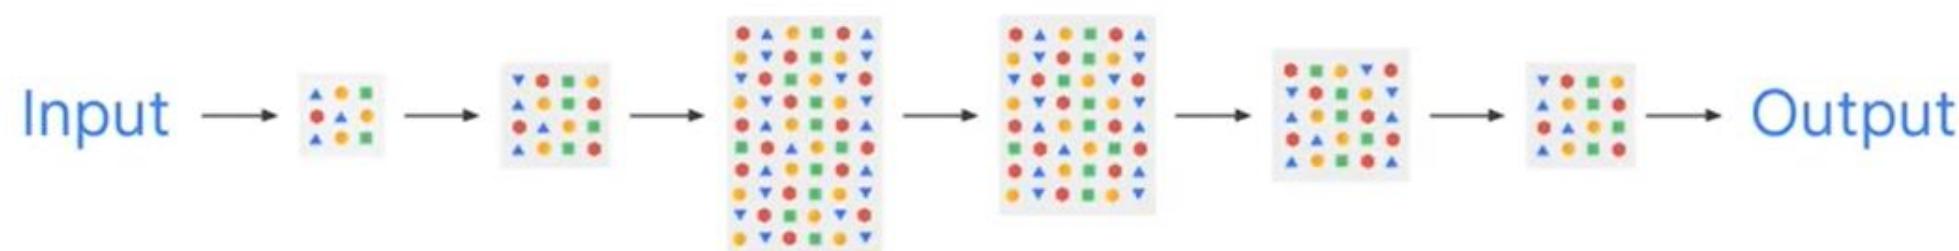
Stage 2: Predict with a trained model





**Focus on both the
training and
inference stages
of ML**

Neural networks are one important technology we use



Deep learning has come a long way in just the past few years



Google Photos

illustrates how far ML has come.



Google Translate

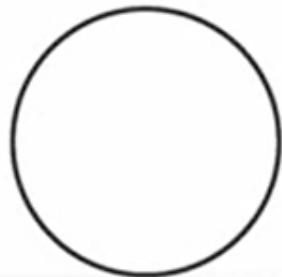
is a combination of several models.



Gmail

Smart Reply Inbox
20% of all responses sent on mobile.

Google Search, our flagship application



The image shows a Google search interface. The search bar at the top contains the query "giants". Below the search bar, there are several search results:

- New York Giants**
Football team
- giants vs dodgers**
- giants schedule**
- San Francisco Giants**
Baseball team

Machine learning scales better than hand-coded rules



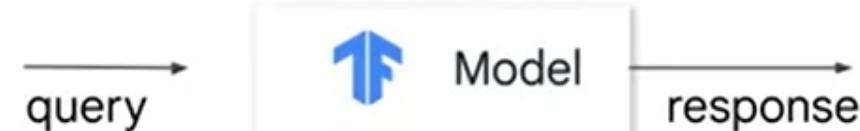
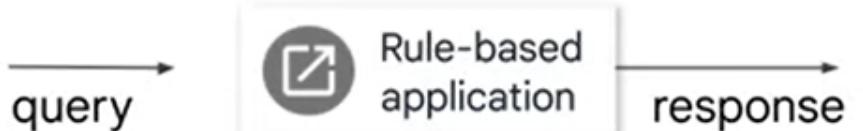
RankBrain (a deep neural network for search ranking)
improved performance significantly

#3 Signal for search ranking, out of hundreds

#1 Improvement to ranking quality in 2+ years

Google thinks of ML as the way
to **scale**, to **automate**, to **personalize**

ML can be used to solve many problems for which you are writing rules today



-
- Code up rules based on human expertise
 - Apply rules program to make decisions
 - Add new rules in response to bug reports

-
- Train model based on data
 - Deploy model at scale to make predictions
 - Continuously train model on data

ML converts examples into knowledge

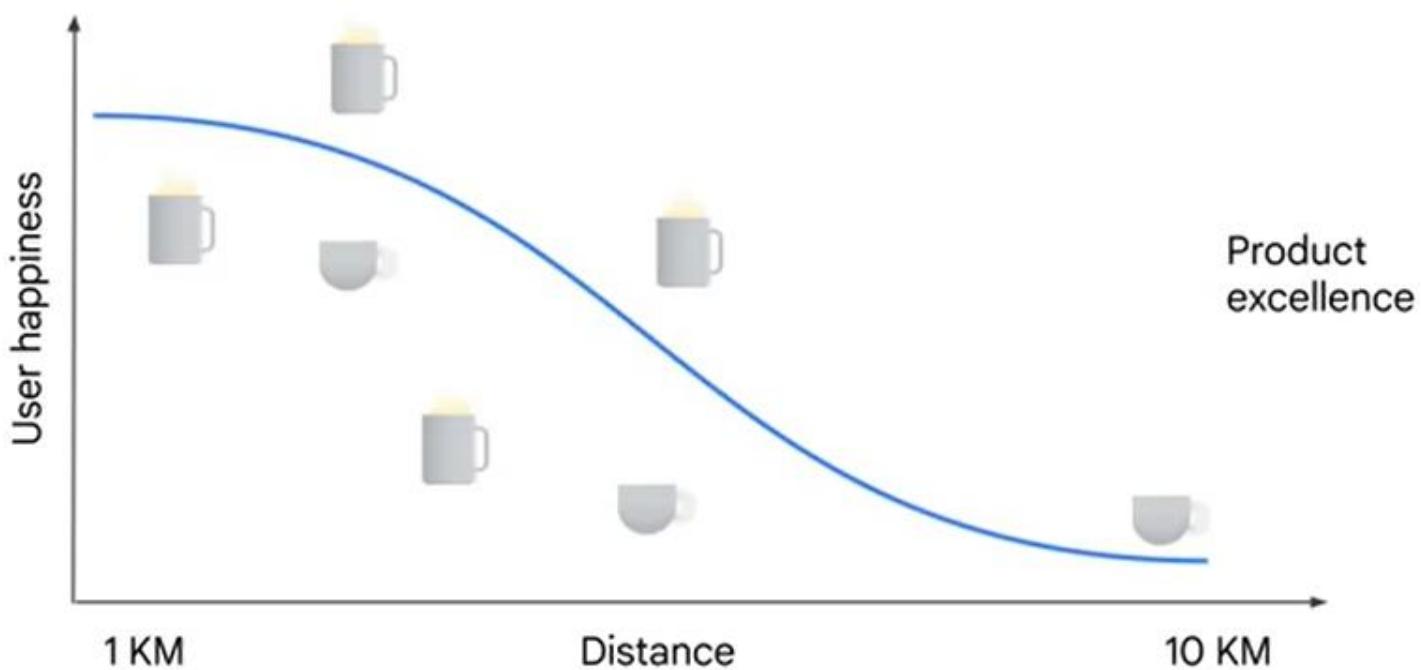


Google

Coffee near me



Good learning involves blending all the users' preferences



Cloud machine learning use cases



Manufacturing

- Predictive maintenance or condition monitoring
- Warranty reserve estimation
- Propensity to buy
- Demand forecasting
- Process optimization
- Telematics



Retail

- Predictive inventory planning
- Recommendation engines
- Upsell and cross-channel marketing
- Market segmentation and targeting
- Customer ROI and lifetime value



Healthcare and Life Sciences

- Alerts and diagnostics from real-time patient data
- Disease identification and risk satisfaction
- Patient triage optimization
- Proactive health management
- Healthcare provider sentiment analysis

Cloud machine learning use cases (Continued)



Travel and Hospitality

- Aircraft scheduling
- Dynamic pricing
- Social media – consumer feedback and interaction analysis
- Customer complaint resolution
- Traffic patterns and congestion management



Financial Services

- Risk analytics and regulation
- Customer segmentation
- Cross-selling and up-selling
- Sales and marketing campaign management
- Credit worthiness evaluation



Energy, Feedstock and Utilities

- Power usage analytics
- Seismic data processing
- Carbon emissions and trading
- Customer-specific pricing
- Smart grid management
- Energy demand and supply optimization

Example solution: Demand forecasting in manufacturing



ML problem:

What is being predicted?

How many units of widget X should you manufacture this month?

What data is needed?

Historical data on # of units sold, price it was sold at, # of units returned, price of competitor product, # of units of all items that use widget X that were sold (e.g. if widget is a phone display panel, how many smartphones were sold, regardless of which display panel they carried?), economic figures (e.g. customer confidence, interest rate), this-month-last-year

Example solution: As a software problem



`predictDemand(widgetID,
month=CurrentTime.month)`

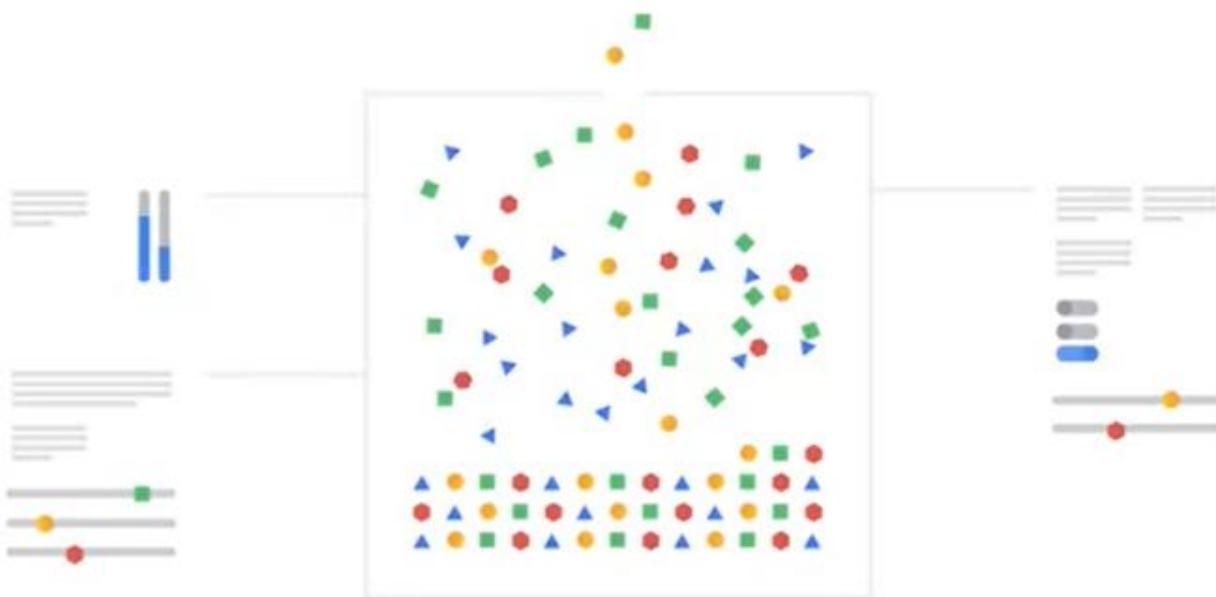
[Who will use this service?](#)

Product managers, logistics
managers

[How are they doing it today?](#)

They examine trends of e.g.
phone sales, overall economy,
trade publications and make a
decision

Example solution: As a data problem



Data problem:

Collect: economic data,
competitor data, industry data,
our figures

Analyze: craft features that our
experts are looking at today from
this data and use as inputs to
model

React: automatic?

Infuse your apps with ML



How much is this
car worth?



1st Guess

TOYOTA
Land Cruiser PRADO

CBA-TRJ150W

89%

Price range (USD)
39,390~43,320

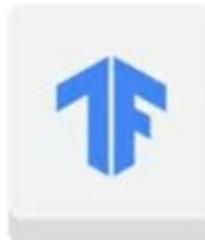


There are pre-trained machine learning services available on Google Cloud

Custom ML models

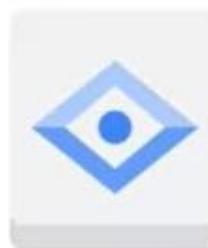


Vertex AI

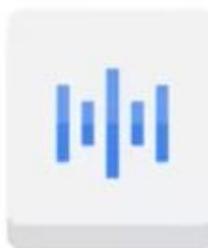


TensorFlow

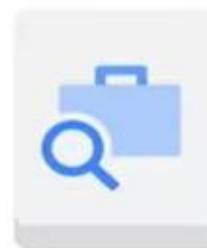
Pre-trained ML Models



Vision API



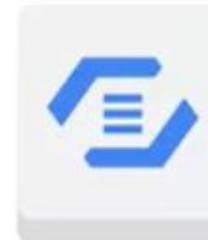
Speech API



Jobs API



Translation API



Natural Language API



Video Intelligence API

Ocado routes emails based on NLP

Improves natural language processing
of customer service claims

"Hi Ocado,
I love your website. I have children so
it's easier for me to do the shopping
online. Many thanks for saving my time!
Regards"

Feedback

Customer is happy



"Thanks to the Google
Cloud Platform, Ocado
was able to use the power
of cloud computing and
train our models in parallel."

The ML marketplace is moving towards increasing levels of ML abstraction



Create a custom image model to price cars.



Build off NLP API to route customer emails.



Use Vision API as-is to find text in memes.



Use Dialogflow to create a new shopping experience.

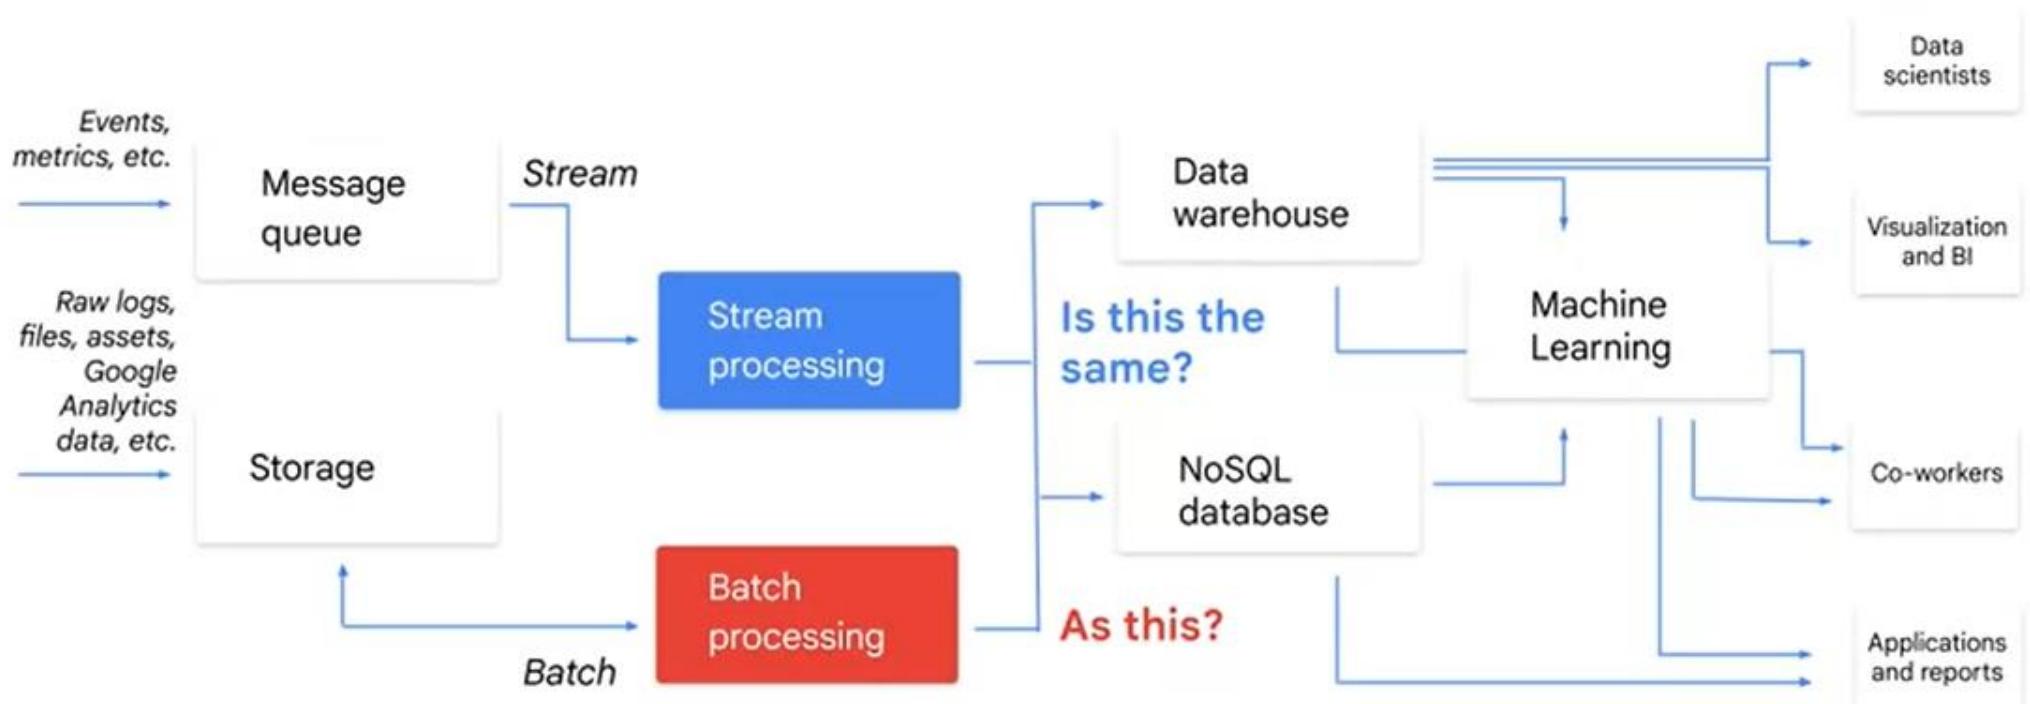
Build a data strategy around ML

If ML is a rocket engine,
data is the fuel.

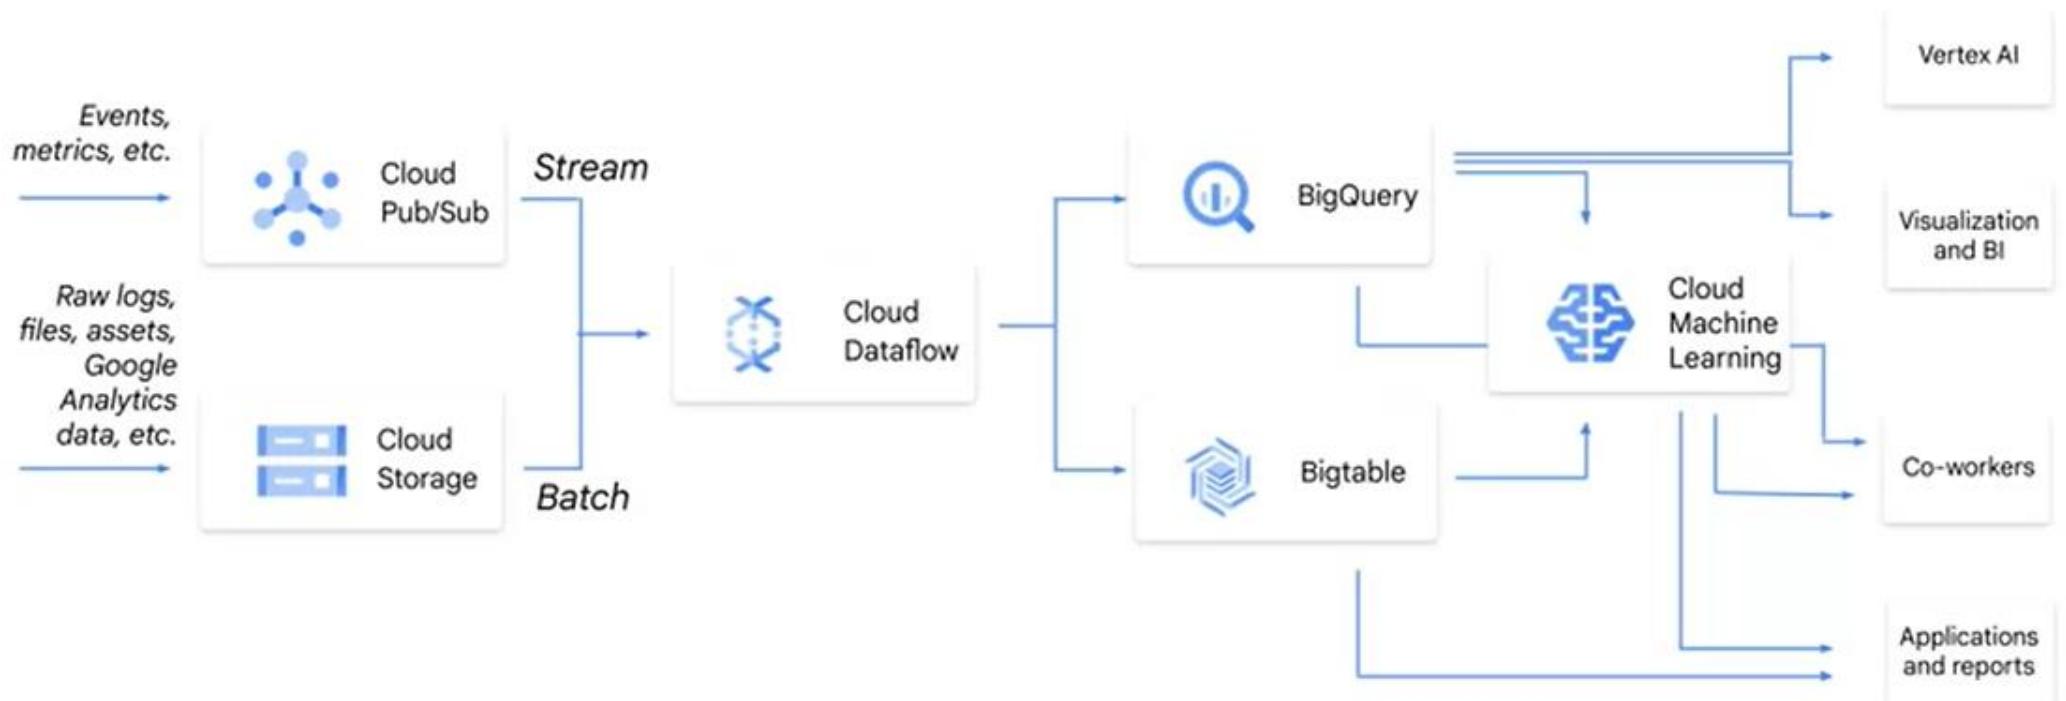


- 1 Collecting data is often the longest and hardest part of an ML project, and the one most likely to fail
- 2 Manual analysis helps you fail fast and try new ideas in a more agile way
- 3 To build a good ML model, you have to know your data
- 4 ML is a journey towards automation and scale

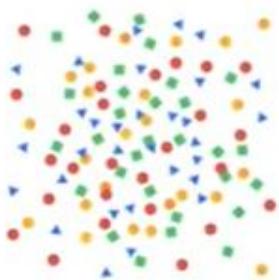
For machine learning, you need to build a streaming pipeline in addition to a batch pipeline



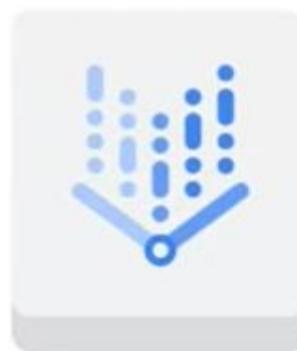
Sophistication around real-time data is key



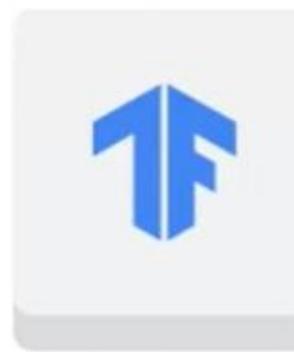
Performance metrics for training are different than for predictions



Training should scale to handle a lot of data.



Vertex AI



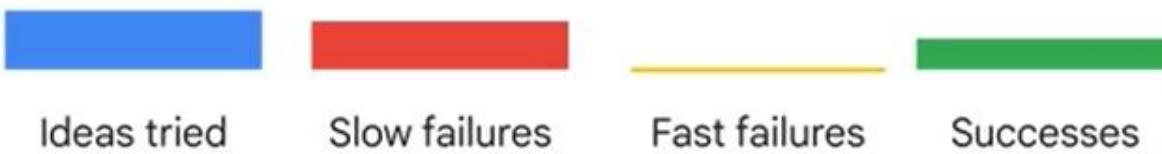
TensorFlow



Predictions should scale to handle large number of queries per second.

Freedom to experiment (and fail) is important

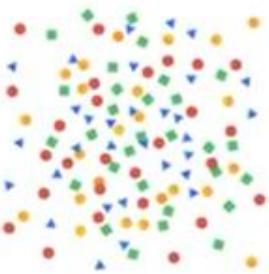
Take your time
and succeed.



Fail fast
and iterate.



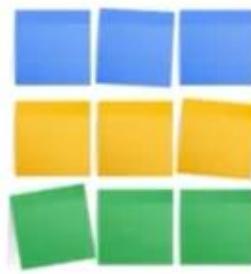
Build on top of Google



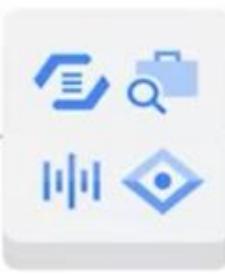
Images
Audio
Video
Free-form text



ML API



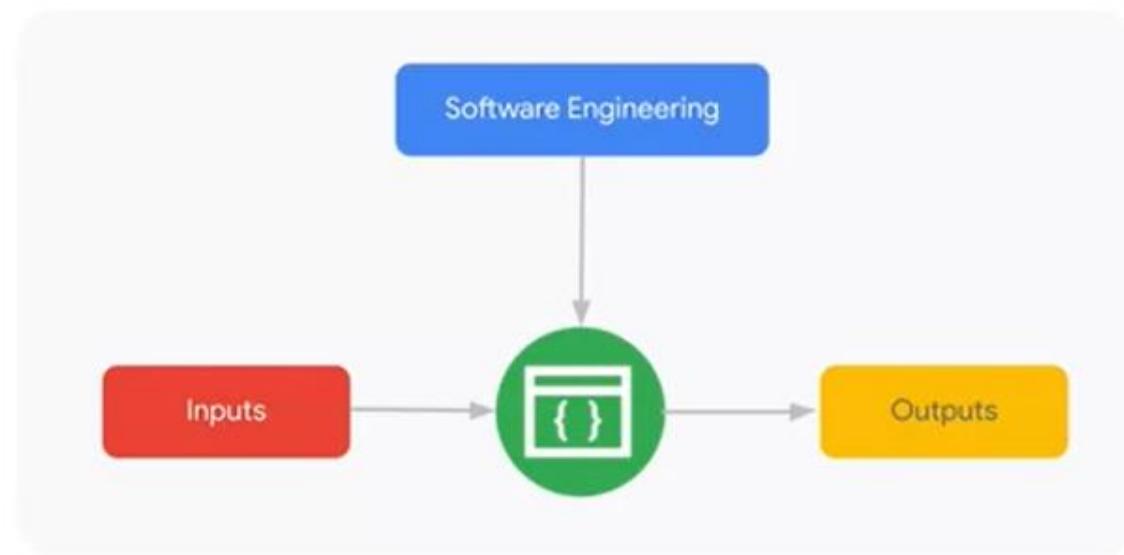
Places
Labels
People
Events
...



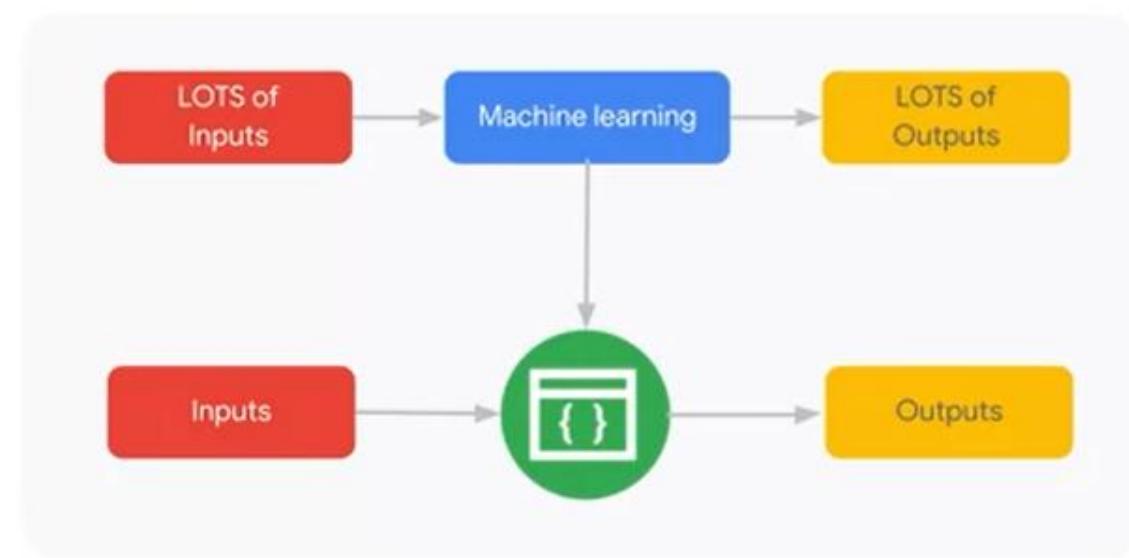
Vertex AI

- Module 1: What it Means to be AI First

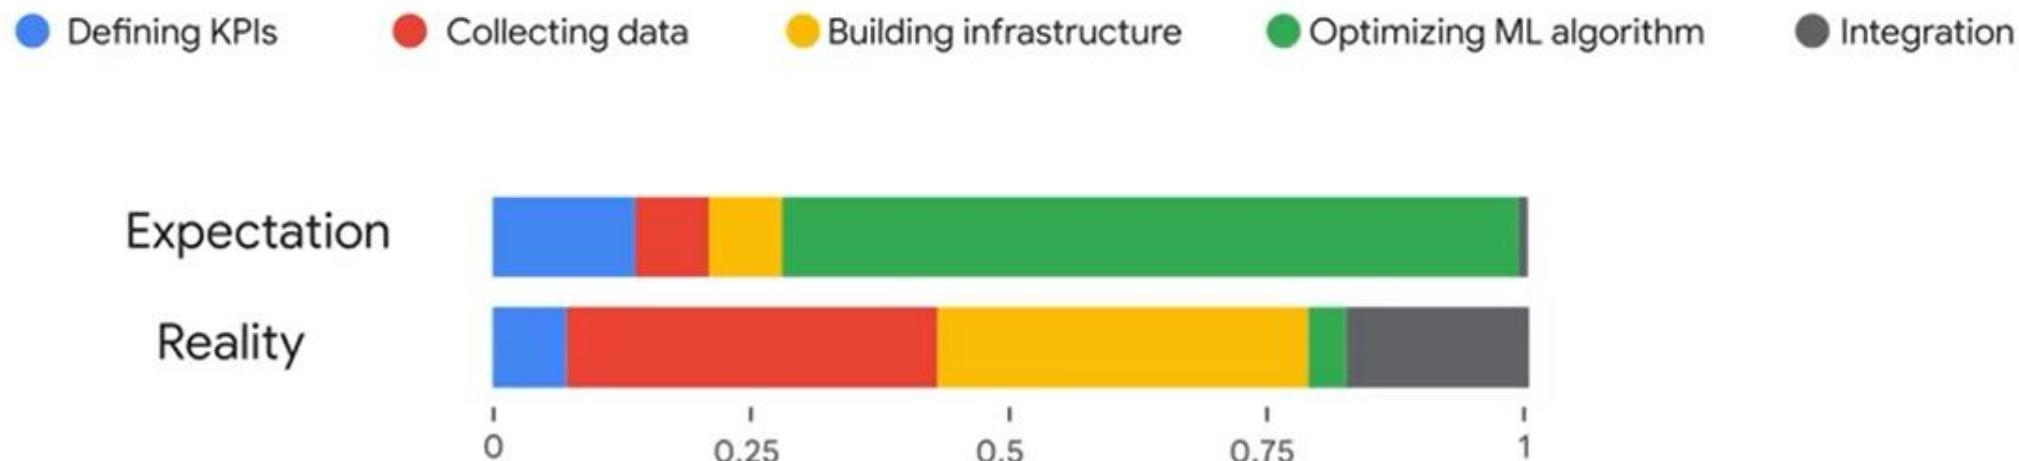
**Software
engineers write
program rules**



**Machine learning
figures out
program rules**



ML effort allocation



Avoid these top 10 ML pitfalls

● Defining KPIs ● Collecting data ● Integration ● Infrastructure ● Optimizing ML

- ● ● 01 ML requires just as much software infrastructure
- 02 No data collected yet
- 03 Assume the data is ready for use
- 04 Keep humans in the loop
- 05 Product launch focused on the ML algorithm
- 06 ML optimizing for the wrong thing
- 07 Is your ML improving things in the real world
- ● 08 Using a pre-trained ML algorithm vs building your own
- 09 ML algorithms are trained more than once
- 10 Trying to design your own perception or NLP algorithm

Ugh, so that's the bad news, what's the good news?

Most ML value
comes along
the way

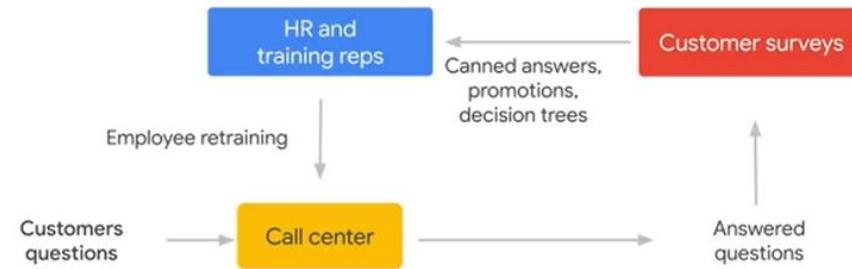
ML improves
almost everything
it touches

If ML is hard, it's
hard for your
competitors too

ML is a great
differentiator

ML and business processes

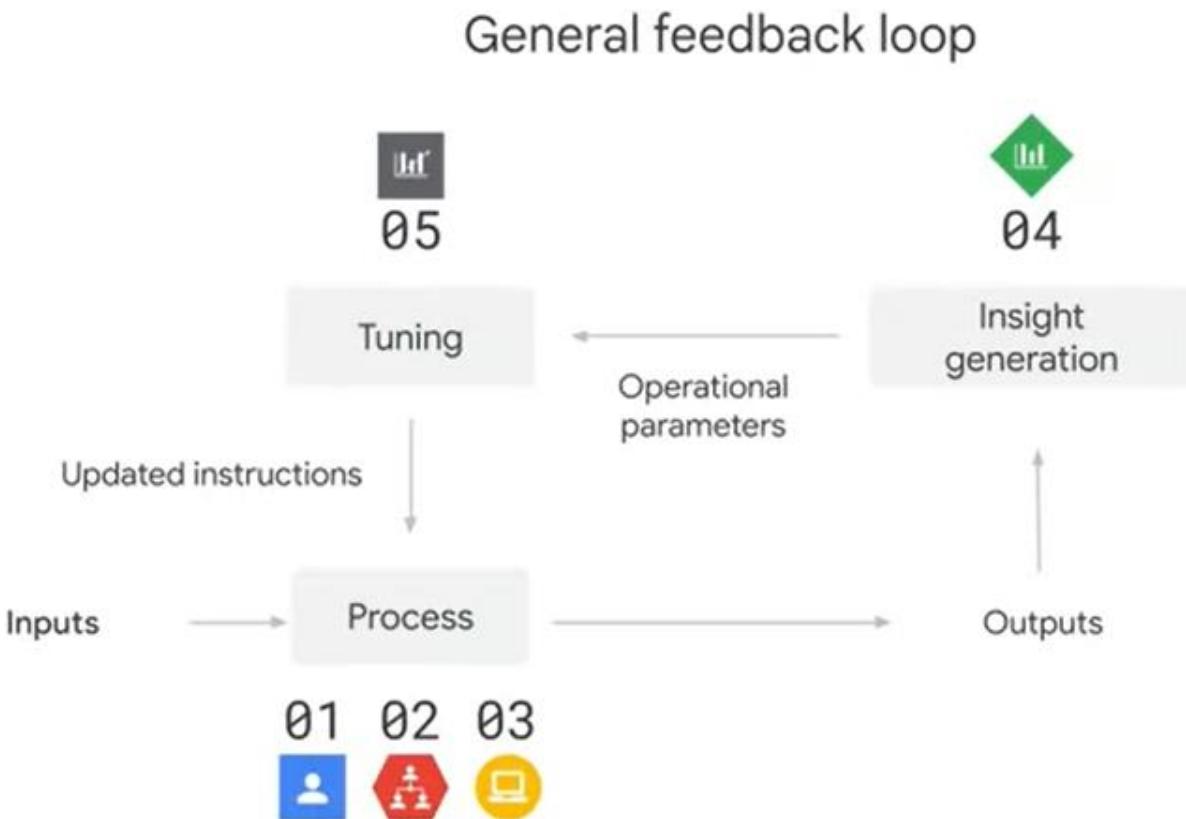
Example: Call center feedback loop



Path to ML: The 5 phases

How change happens in phases:

- 01 Individual contributor
- 02 Delegation
- 03 Digitization
- 04 Big data and analytics
- 05 Machine learning



Prototype and try out ideas



1

Individual
contributor

Dangers of skipping this step:

- Inability to scale.
- Product heads make big, incorrect assumptions that are hard to change later.

Dangers of lingering too long here:

- One person gets skilled and then leaves.
- Fail to scale up the process to meet demand in time.

Gently ramp up to include more people



2

Delegation

Dangers of skipping this step:

- Not forced to formalize the process.
- Inherent diversity in human responses become a testbed--great product learning opportunity.
- Great ML systems will need humans in the loop.

Dangers of lingering too long here:

- Paying a high marginal cost to serve each user.
- More voices will say automation isn't possible.
- Organizational lock-in.

Automate mundane parts of the process



3

Digitization

Dangers of skipping this step:

- You will always need infrastructure.
- IT project and ML success tied and the whole project will fail if either does.

Dangers of lingering too long here:

- Your competitors are collecting data and tuning their offers from these new insights.

Measure and achieve data-driven success



4

Big data and
analytics

Dangers of skipping this step:

- Unclean data means no ML training.
- You can't measure success.

Dangers of lingering too long here:

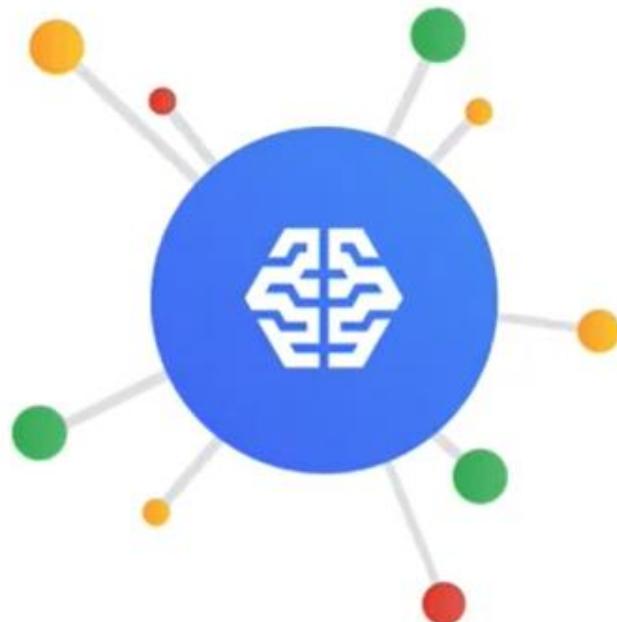
- Limit the complexity of problems you can solve.

Automated feedback loop that can outpace human scale

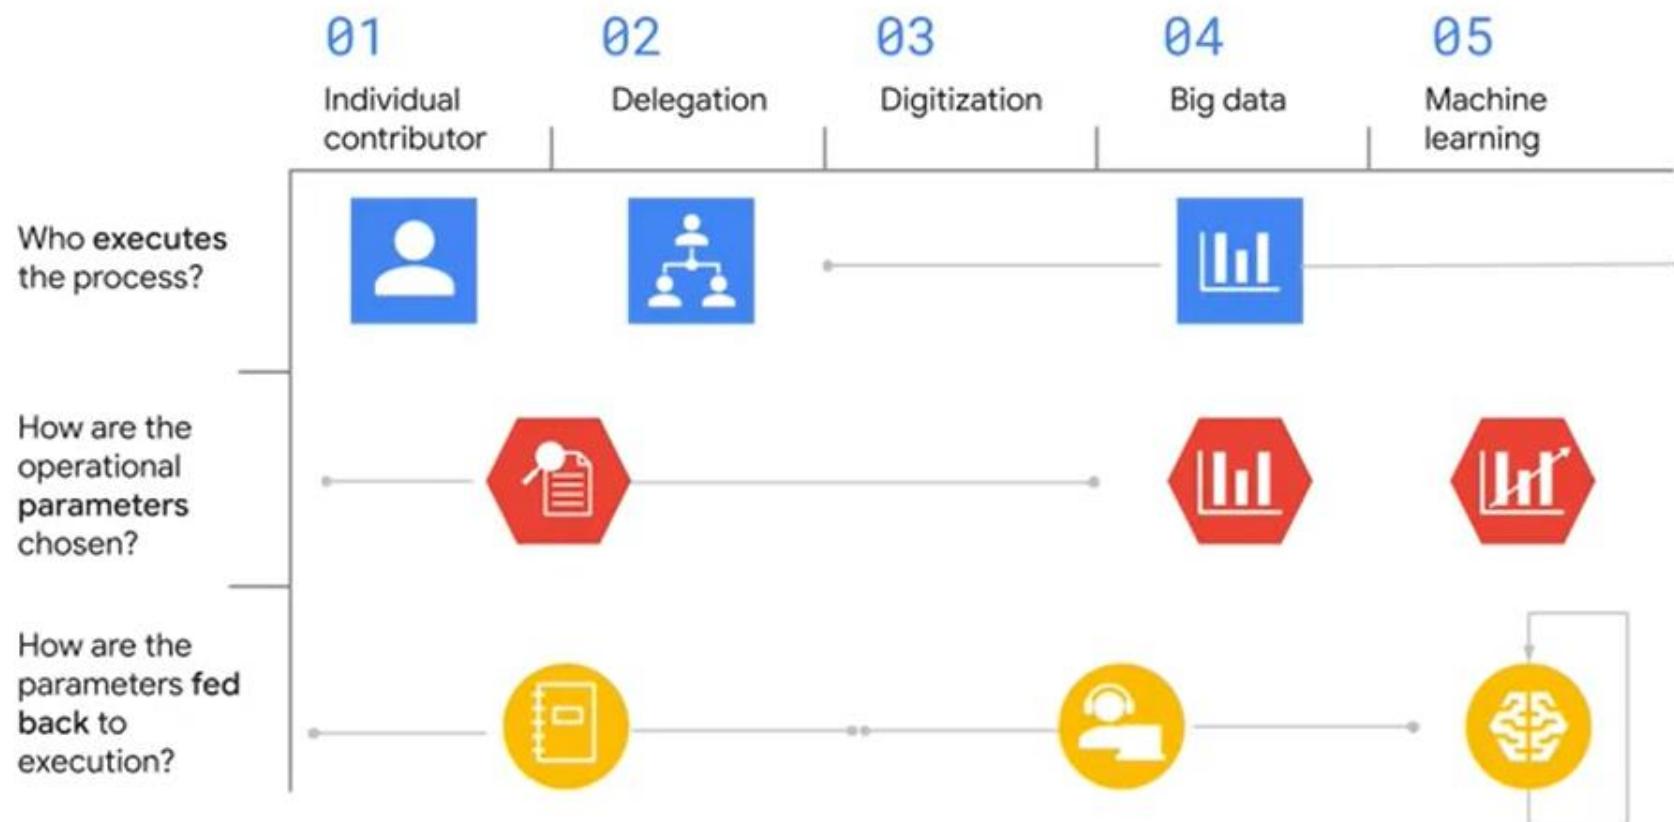


5

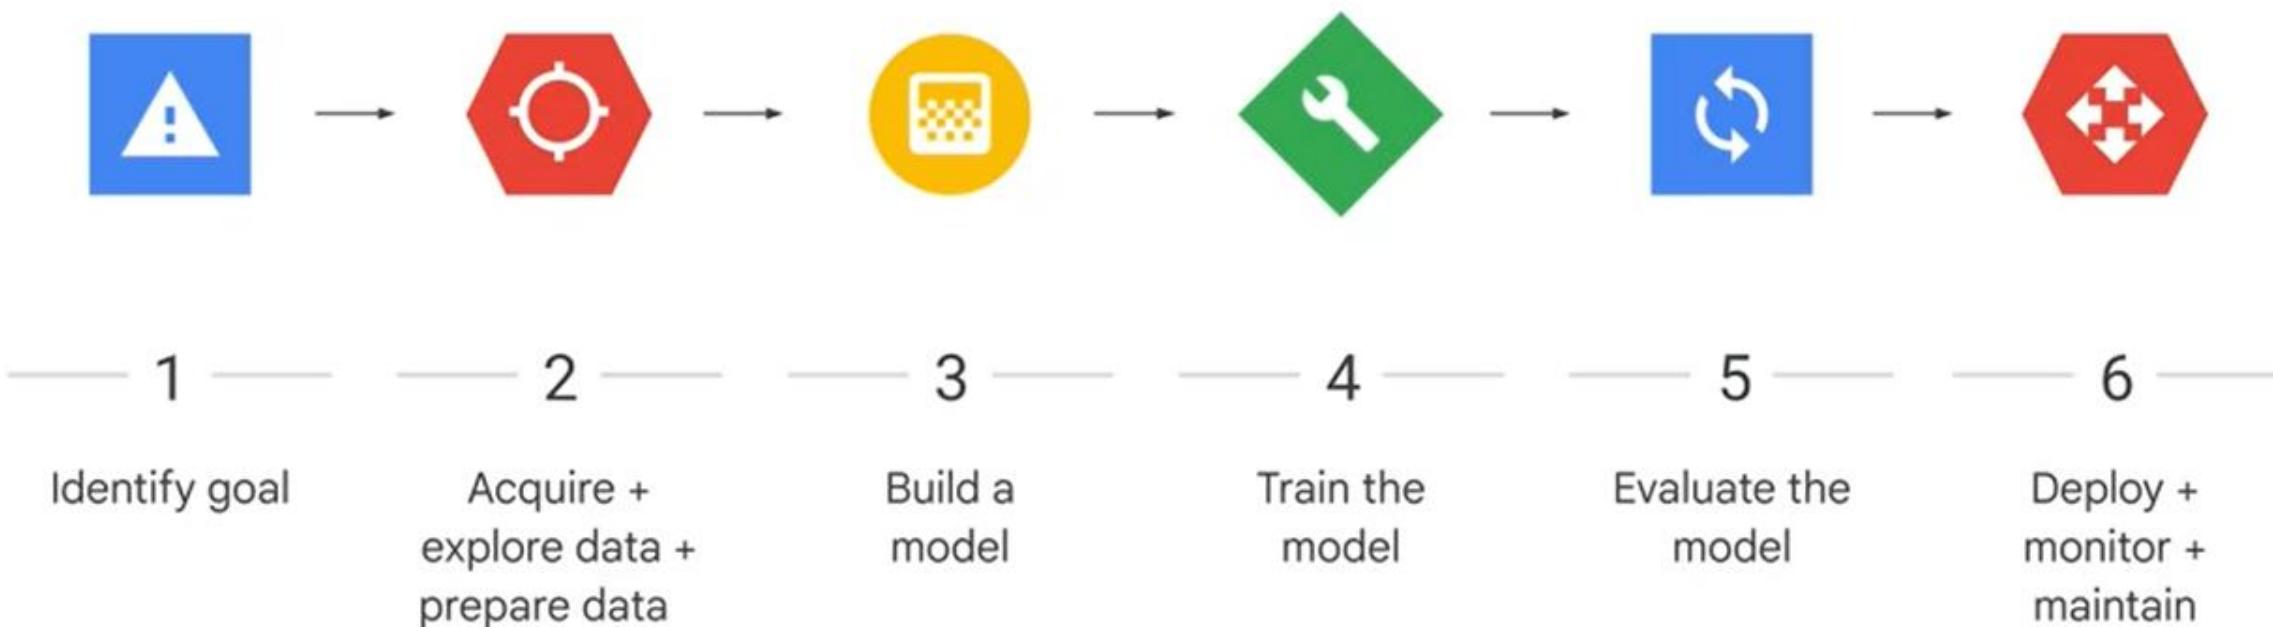
Machine
learning



Reviewing the path to ML: 5 phases



To build a machine learning model for production



Typical ML development during experimentation

Framing the problem

Prepare training data

Experimenting

Evaluating the model

- Model A
- Model B
- Model C

Typical ML development during experimentation

Different architectures	Different input data sets	Different hyperparameters	Different hardware
<ul style="list-style-type: none">• CNN• RNN• Sorting/Clustering• GANs			

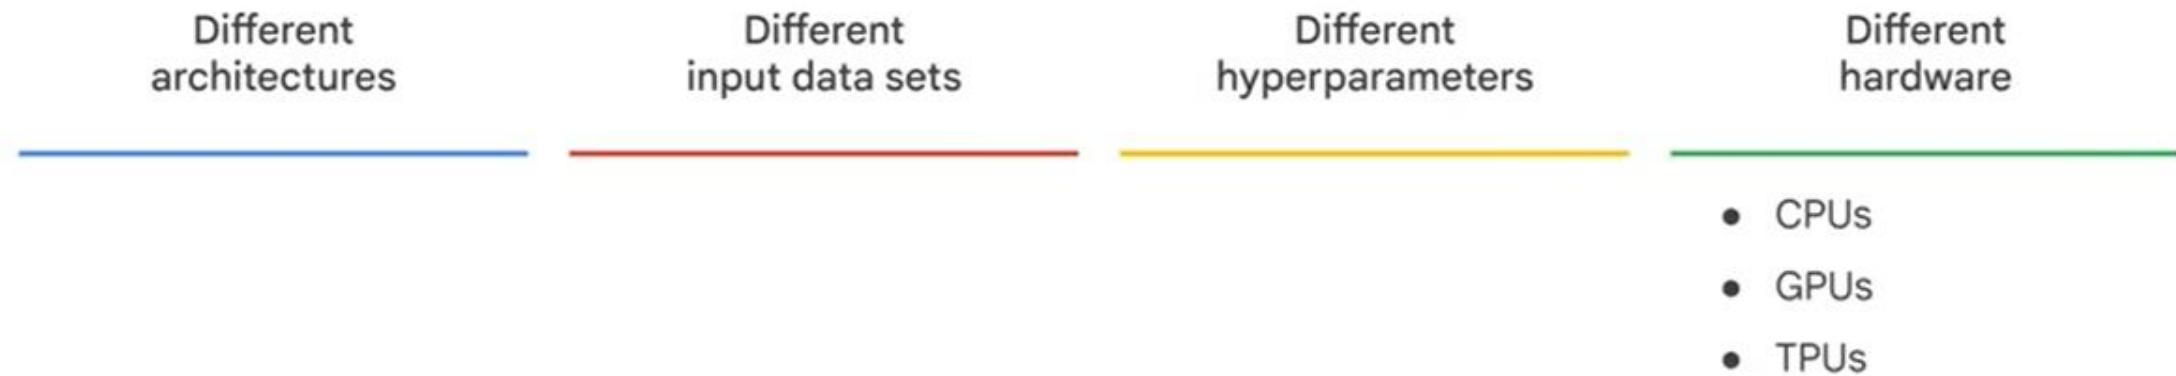
Typical ML development during experimentation

Different architectures	Different input data sets	Different hyperparameters	Different hardware
	<ul style="list-style-type: none">• Numerical data sets• Bivariate data sets• Multivariate data sets• Categorical data sets• Correlation data sets		

Typical ML development during experimentation

Different architectures	Different input data sets	Different hyperparameters	Different hardware
		<ul style="list-style-type: none">• Learning rate• Number of layers• Num_estimators• Max_depth	

Typical ML development during experimentation



ML application
generates a REST
service for use by a
medical application

Medical application

Baby weight predictor

Example application to predict a baby's weight.

Mother's age 27

Gestation weeks 38

Plurality Single

Baby's gender Male Female Unknown

PREDICT

Prediction 7.19 lbs.



Request

Example:
-Age
-Gestation
-Weeks
-Gender

Prediction

Example:
-Baby's weight

ML application (or its pipeline)

Why model monitoring

- Stale model - The underlying data distribution has shifted over time.
- Misconfigured model in production deployment.

Baby weight predictor

Example application to predict a baby's weight.

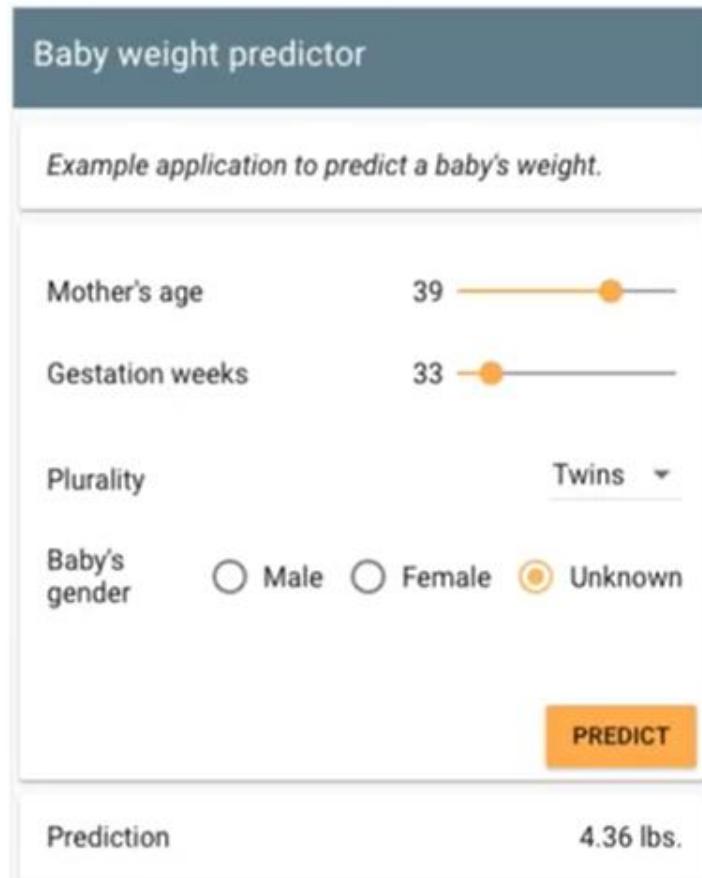
Mother's age 39

Gestation weeks 33

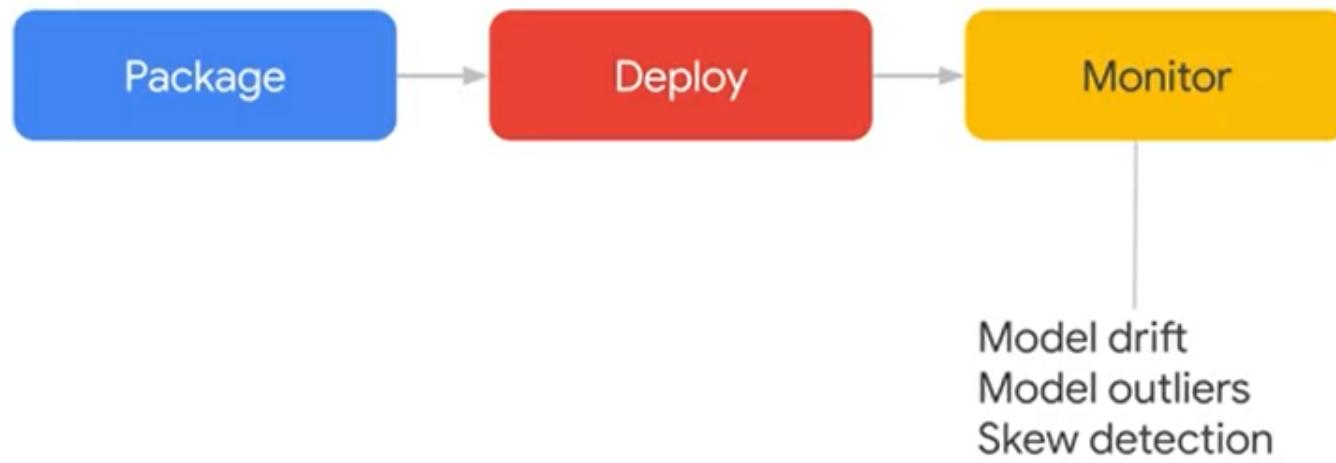
Plurality Twins

Baby's gender Male Female Unknown

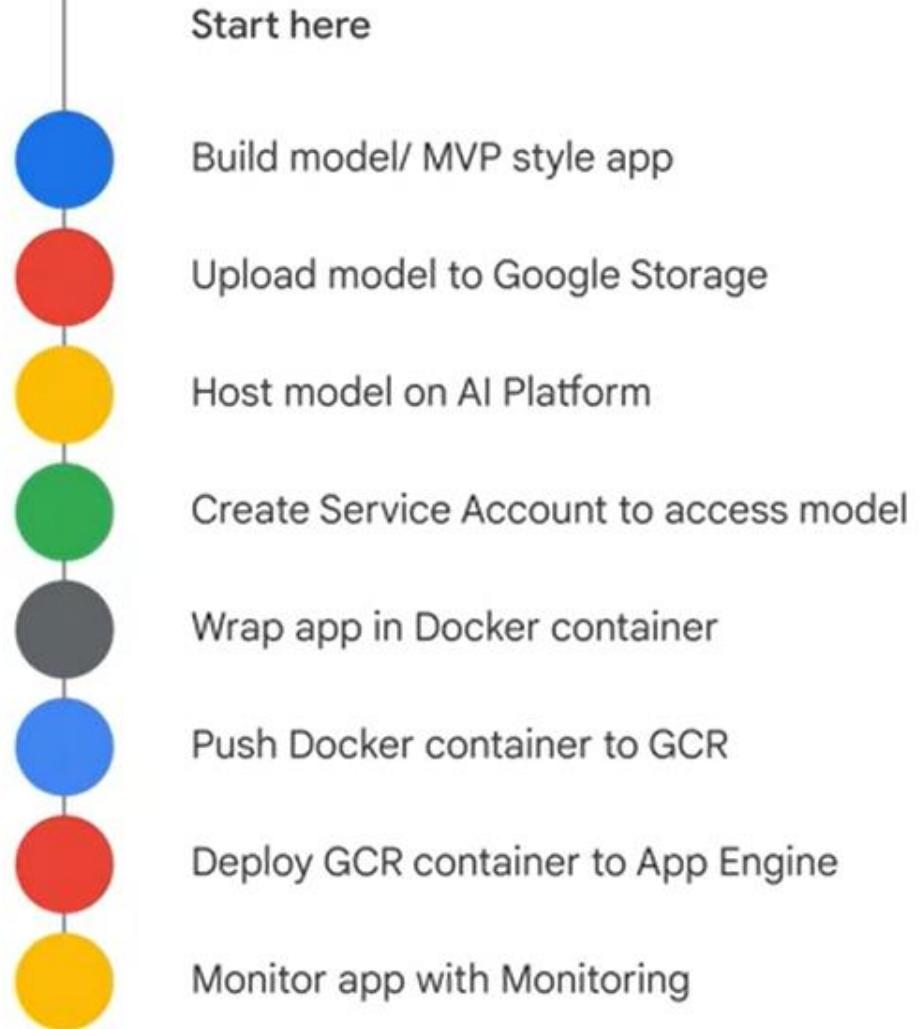
Prediction 4.36 lbs.



Moving from
experimentation to
production requires
packaging, deploying,
and monitoring
your model



ML product knowledge required



What is there to unify?



Dataset is

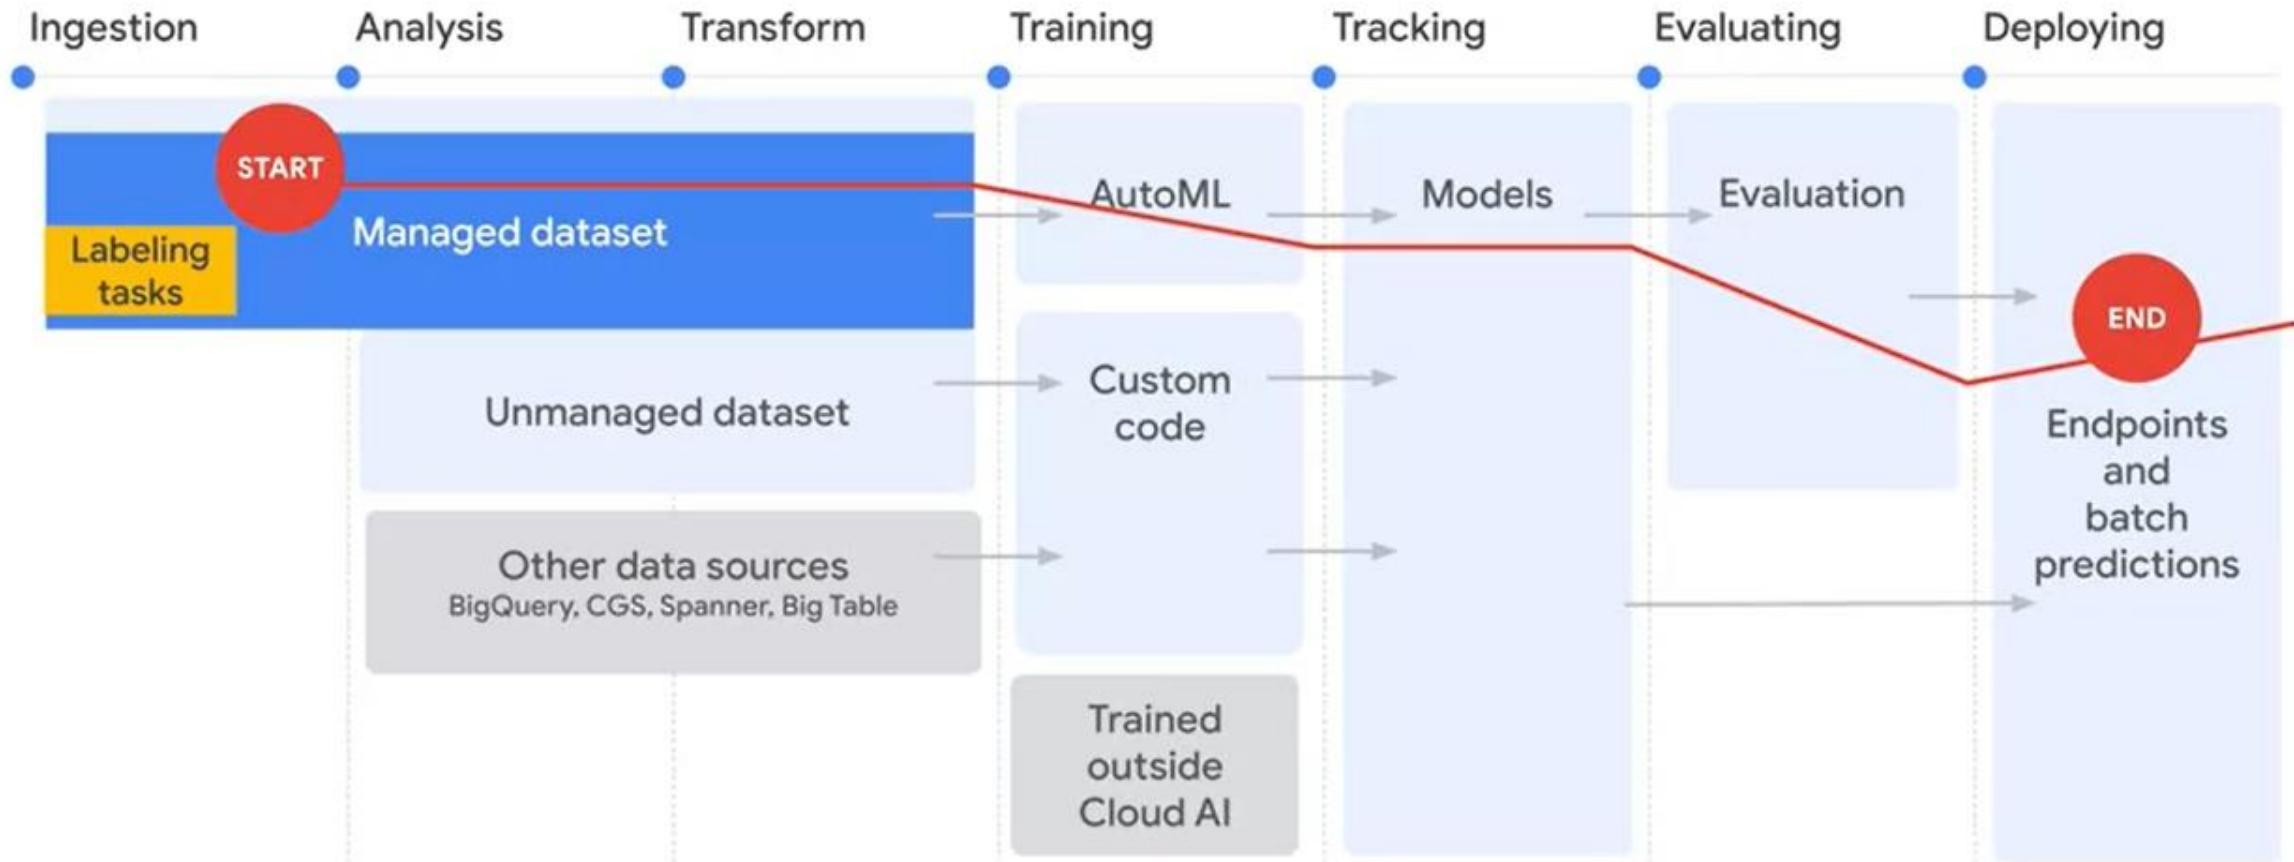
- Created
- Ingested
- Analyzed
- Cleaned (ETL or ELT)



Model is

- Trained, which includes experimentation and hypothesis testing, and hyperparameter tuning.
- Versioned and rebuilt when there is new data, on a schedule, or when the code changes (ML Ops).
- Evaluated and compared to existing model versions.
- Deployed and used for online and batch predictions.

Vertex AI



Vertex AI provides a unified set of APIs for the ML lifecycle. Diagram courtesy Henry Tappen and Brian Kobashikawa

Choose a training method

AutoML

- Create and train a model with minimal technical effort.
- Quickly prototype models or explore datasets before developing in a custom training application.

Custom training

- Create a training application optimized for your targeted outcome.
- Maintain complete control over training application functionality.
 - Target any objective, use any algorithms, develop your own loss functions or metrics, or make other customizations.

Vertex AI



Fast experimentation



Accelerated deployment



Simplified model
management

Components of Vertex AI

Vertex AI > Workbench

The screenshot shows the Vertex AI Workbench dashboard. On the left, a sidebar menu lists various components: Dashboard, Datasets, Features, Labeling tasks, **Workbench** (which is highlighted with a red box), Pipelines, Training, Experiments, Models, Endpoints, Batch predictions, and Marketplace. The main area is titled "Dashboard" and contains a large blue callout box with the text: "Workbench provides a Jupyter notebook development environment for the entire ML workflow. Access data, process data in a Dataproc cluster, train a model, and share results." Below this, there's a "Region" dropdown set to "us-central1 (Iowa)". The dashboard is divided into three main sections: "Prepare your training data", "Train your model", and "Get predictions". Each section has a brief description and a "CREATE" button: "+ CREATE DATASET", "+ TRAIN NEW MODEL", and "+ CREATE BATCH PREDICTION". To the right of these sections is a decorative illustration featuring a lightbulb, clouds, and various data visualization icons like a network graph and a line chart.

Vertex AI > Pipelines

Vertex AI

Dashboard

- Dashboard
- Datasets
- Features
- Labeling tasks
- Workbench
- Pipelines
- Training
- Experiments
- Models
- Endpoints
- Batch predictions
- Marketplace

Get started with Vertex AI

Vertex AI empowers machine learning developers, data scientists, and ML engineers to build, train, and deploy machine learning models at scale.

Pipelines help you to automate, monitor, and govern your ML systems. Each individual part of your pipeline workflow is a component that is defined by code.

Region: us-central1 (Iowa)

Prepare your training data
Collect and prepare your data, then import it into a dataset to train a model
[+ CREATE DATASET](#)

Train your model
Train a best-in-class machine learning model with your dataset. Use Google's AutoML, or bring your own code.
[+ TRAIN NEW MODEL](#)

Get predictions
After you train a model, you can use it to get predictions, either online as an endpoint or through batch requests
[+ CREATE BATCH PREDICTION](#)

2:12 / 5:40

Vertex AI > Metadata

Vertex AI

Dashboard

- Dashboard
- Datasets
- Features
- Labeling tasks
- Workbench
- Pipelines
- Training
- Experiments
- Models
- Endpoints
- Batch predictions
- Metadata

Get started with Vertex AI

Vertex AI empowers machine learning developers, data scientists, and data engineers to take their projects from ideation to deployment, quickly and cost-effectively. [Learn more](#)

ENABLE VERTEX AI API

Region: us-central1 (Iowa)

Vertex ML Metadata stores artifacts and metadata for pipelines run using Vertex AI Pipelines.

CREATE DATASET

Train your model

Train a best-in-class machine learning model with your dataset. Use Google's AutoML, or bring your own code.

+ TRAIN NEW MODEL

Get predictions

After you train a model, you can use it to get predictions, either online as an endpoint or through batch requests

+ CREATE BATCH PREDICTION



Machine learning development with Vertex Notebooks

Vertex AI Workbench provides two Jupyter notebook-based options for your data science workflow

Managed notebooks

User-managed notebooks

Deep Learning VM Images

Good choice for data exploration, analysis, modeling, or as part of an end-to-end data science workflow.

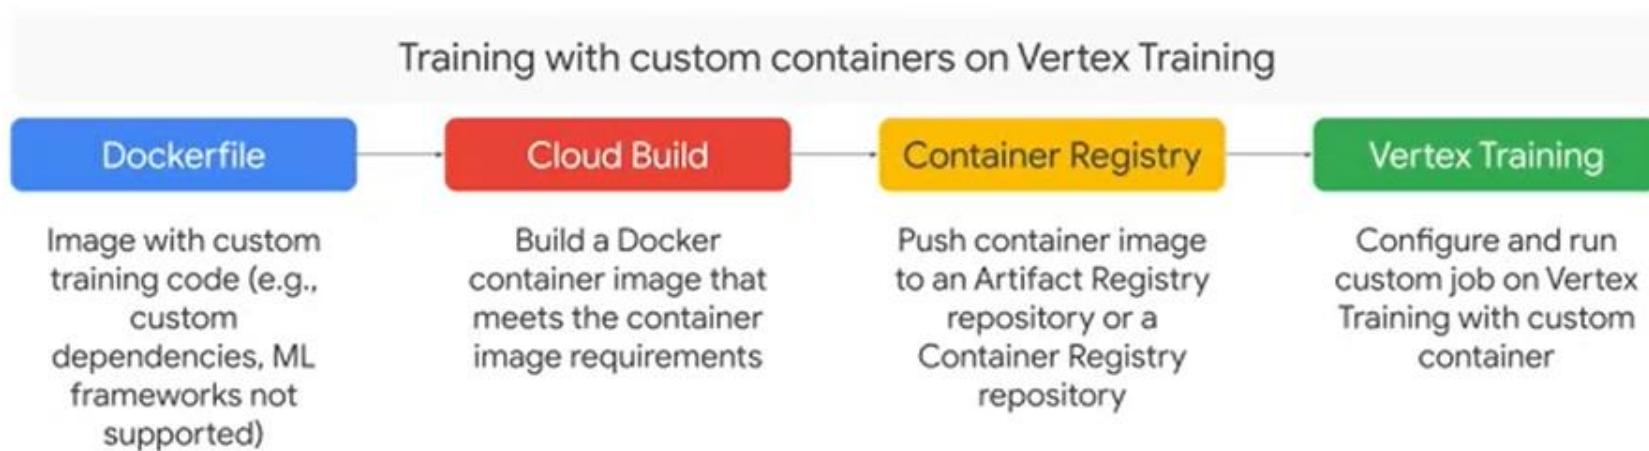
Perform workflow-oriented tasks without leaving the JupyterLab interface.

Managed notebooks

The screenshot shows the Vertex AI Workbench interface with the 'Managed notebooks' section highlighted. On the left, there's a sidebar with icons for Vertex AI, Features, Labeling tasks, Workbench (which is selected), Pipelines, Training, Experiments, Models, and Endpoints. The main area has a blue header 'Managed notebooks' with tabs for 'MANAGED NOTEBOOKS' (selected), 'USER-MANAGED NOTEBOOKS', and 'EXECUTIONS'. Below the tabs, it says 'Managed notebooks provide JupyterLab services and flexible computing resources integrated with Google Cloud services.' A 'Region' dropdown is set to 'us-central1 (Iowa)'. A 'Filter' input field is present. A table lists three managed notebooks:

Notebook name	Location
managed-notebook-1635869666	OPEN JUPYTERLAB us-central1-b
xyz-team-10-14-2021-gs	OPEN JUPYTERLAB us-central1-f

Managed notebooks



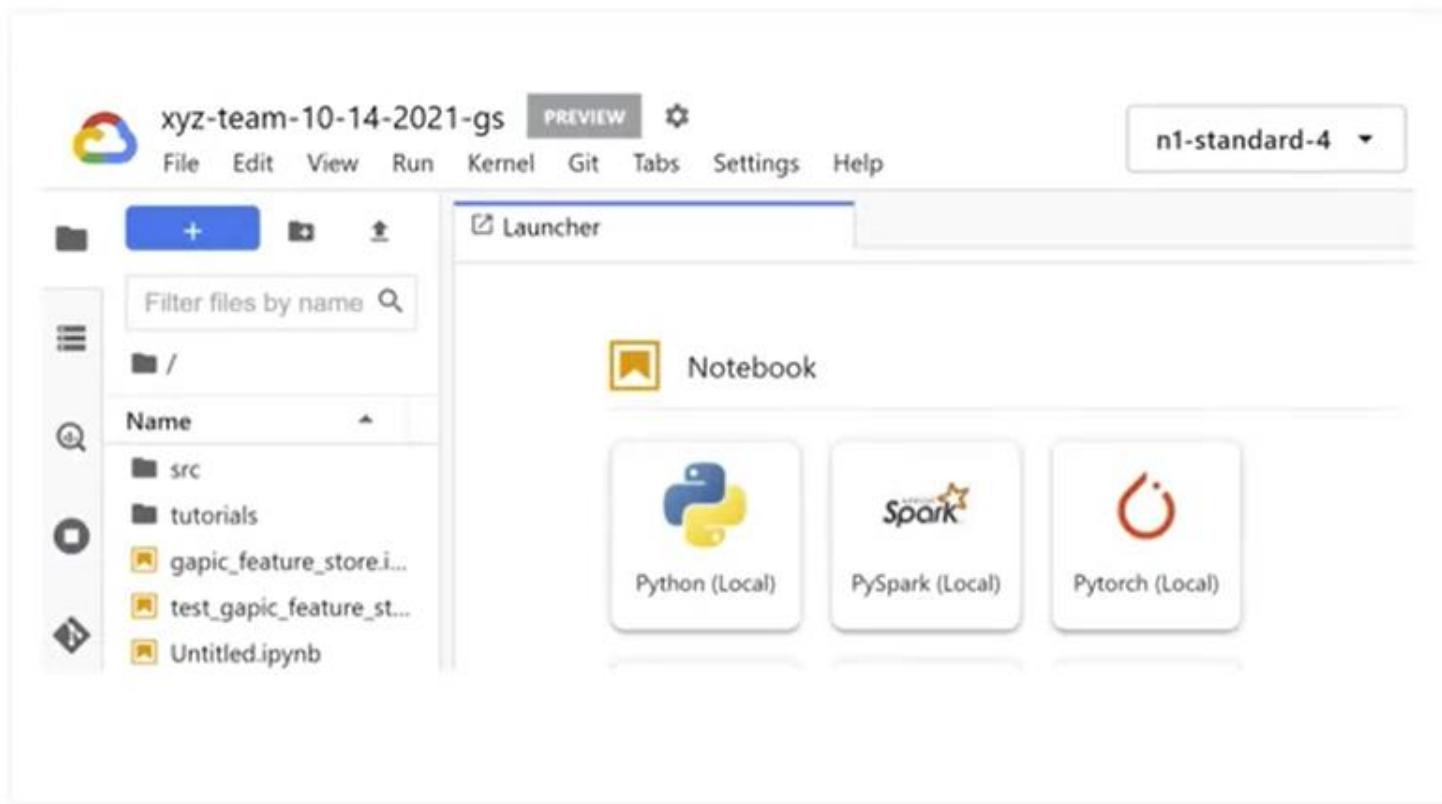
Control hardware

Determine compute resources (GPUs, RAM).

Custom containers

Use TensorFlow, Pytorch, PySpark, R. Add custom Docker container images.

Managed notebooks



Control hardware

Determine compute resources (GPUs, RAM).

Custom containers

Use TensorFlow, Pytorch, PySpark, R. Add custom Docker container images.

Access to data

Use Cloud Storage and BigQuery extension to browse data.

Dataproc integration

Process data quickly by running a notebook on a Dataproc cluster.

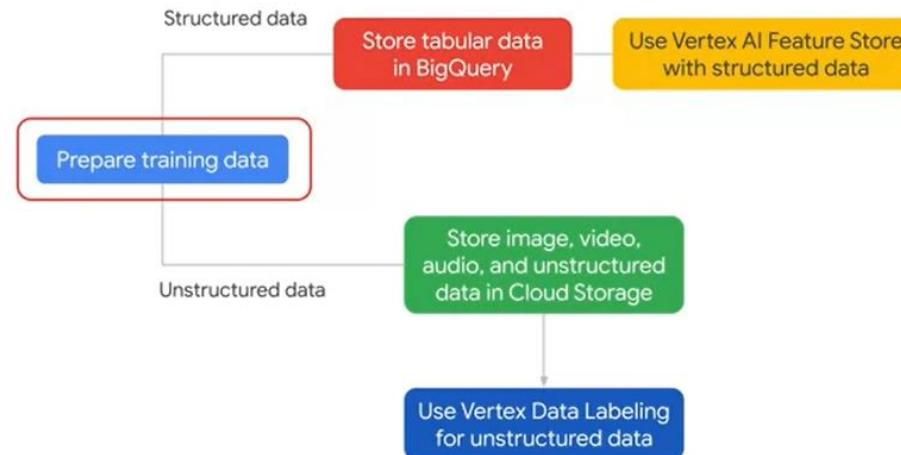
Best practices for machine learning development – Vertex AI

Best practices for preparing and storing data

Data

How it is prepared and stored

Avoid storing data in block storage



Vertex AI Feature Store

Feature Store

Use Feature Store with
structured data

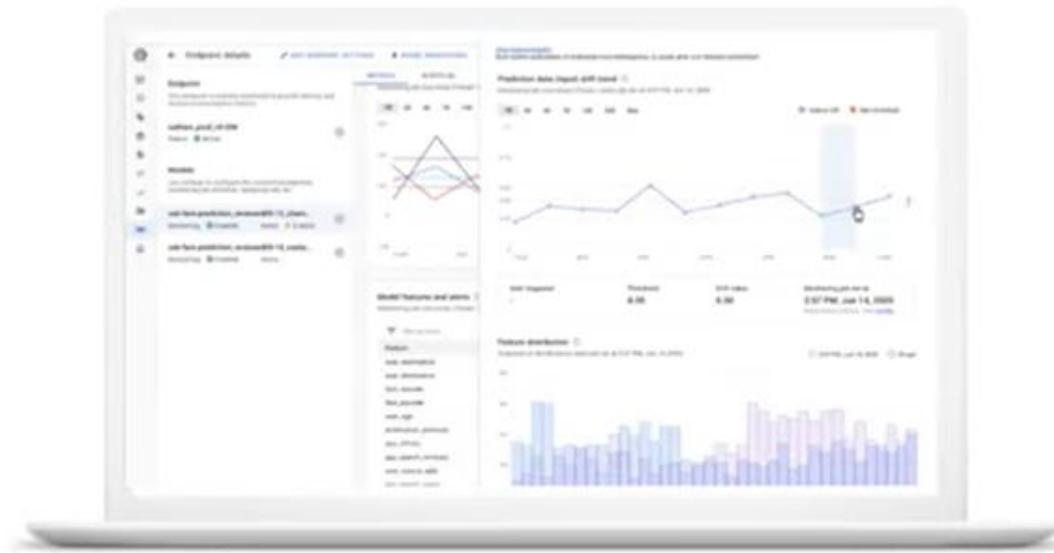
Follow these steps:

1. [Search Vertex AI Feature Store](#)
 - a. Search to see if a feature already exists.
 - b. Fetch those features for your training labels using [Vertex AI Feature Store's batch serving capability](#).
2. [Create a new feature](#)
 - a. Create a new feature using your Cloud Storage bucket or BigQuery location. OR
 - b. Fetch raw data from your data lake and write your scripts to perform feature processing.
 - c. Join the feature values and the new feature values. Merging those feature values produces the training data set.

Best practices for training a model

Model

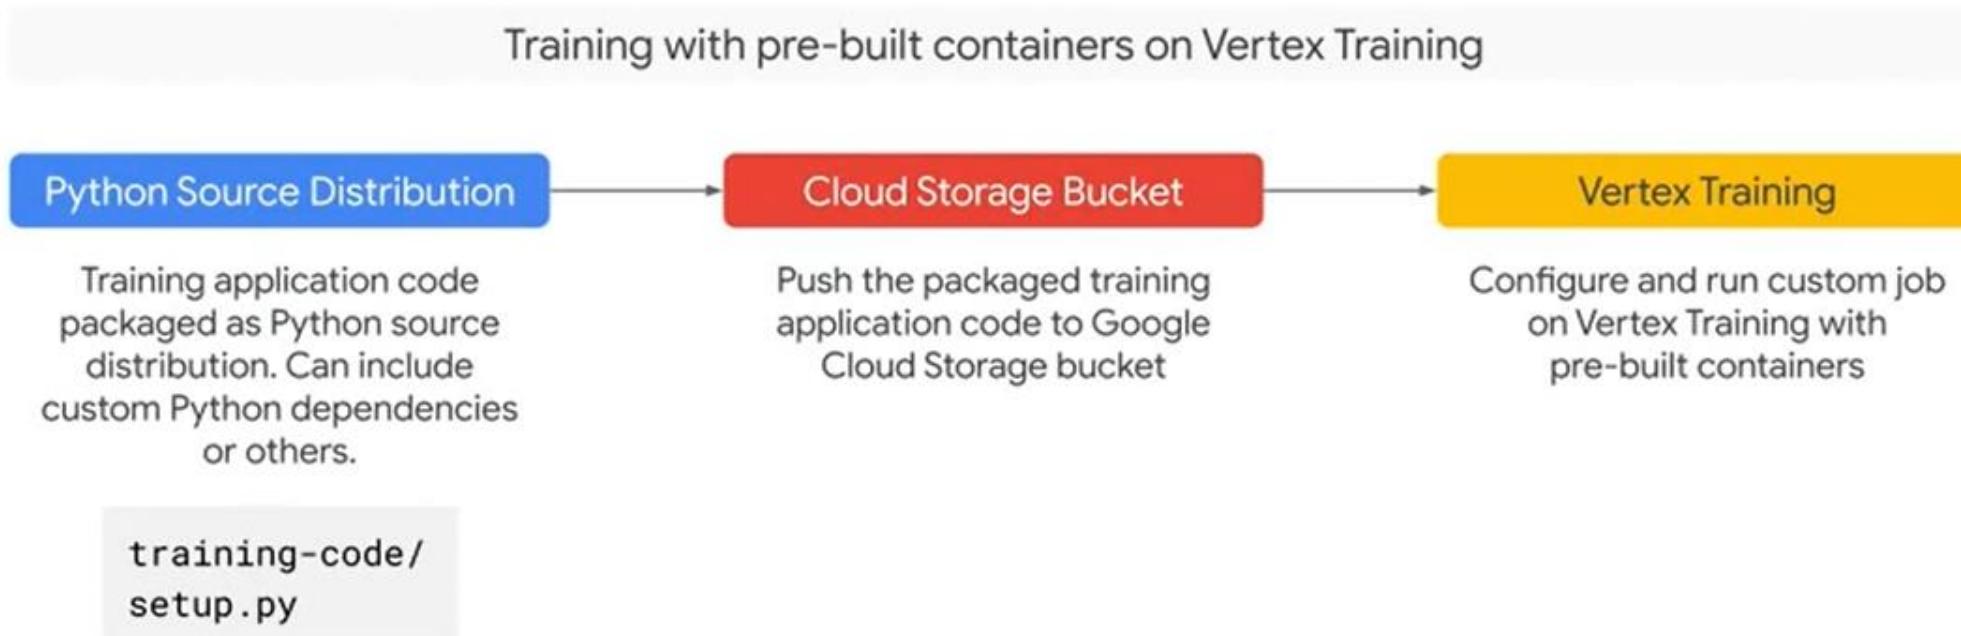
Tips for training, maximizing predictive accuracy, and feature attributions for insights



For small datasets, train a model within the [Notebooks instance](#).

For large datasets, distributed training, or scheduled training, use the [Vertex training service](#).

Training with pre-built containers on Vertex AI



Best practices for Explainable AI



Model

Tips for training, maximizing predictive accuracy, and feature attributions for insights

Offers feature attributions to provide insights into why models generate predictions.

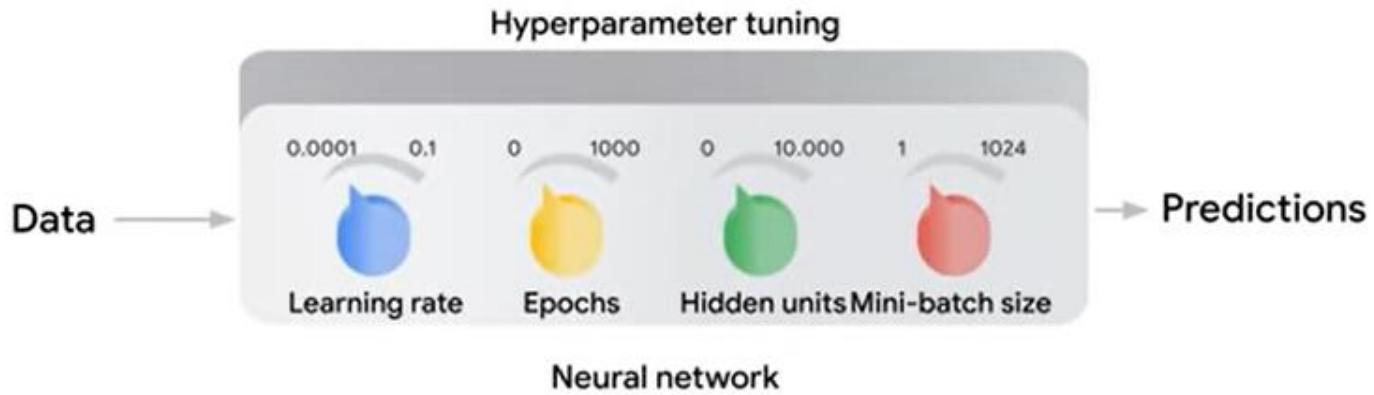
Details the importance of each feature that a model uses as input to make a prediction.

Supports custom-trained models based on tabular and image data.

Hyperparameter tuning with Vertex Training

Model

Maximize your model's predictive accuracy with hyperparameter tuning



The hyperparameters are knobs that act as the network-human interface.

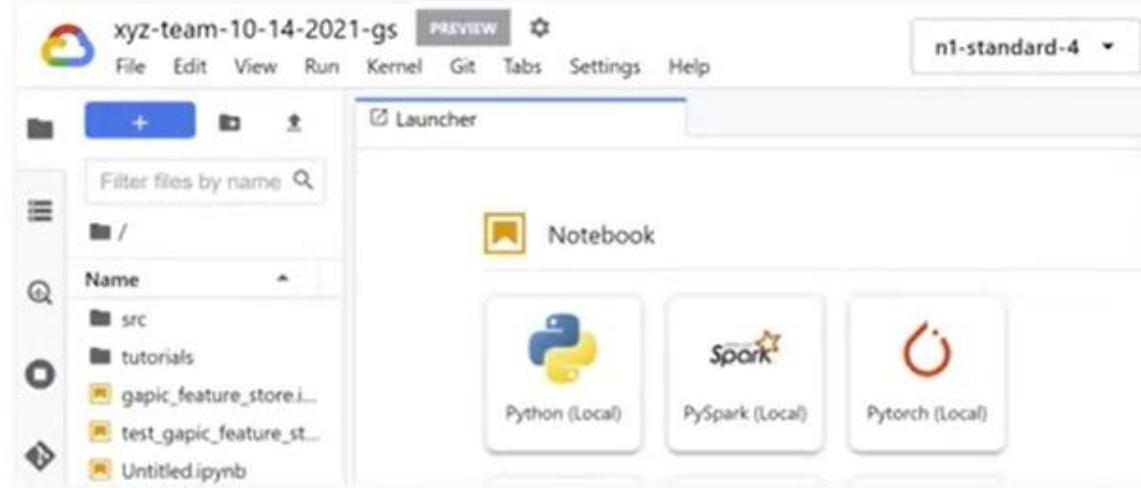
Maximize a model's predictive accuracy. [Vertex Training](#) provides an automated model enhancer to test different hyperparameter configurations when training your model.

No need to manually adjust hyperparameters over the course of numerous training runs to arrive at the optimal values.

Best practices for using Workbench Notebooks

Workbench Notebooks

Use Notebooks to evaluate and understand your models.

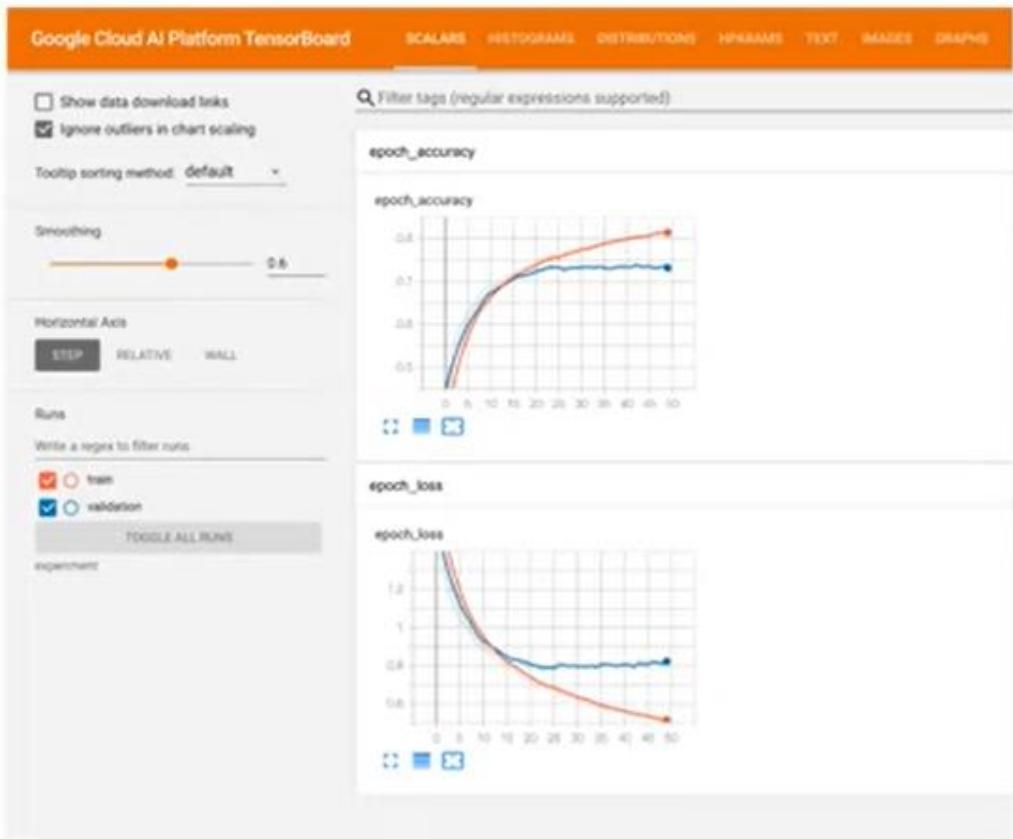


Use [Notebooks](#) to evaluate and understand your models. In addition to built-in common libraries like scikit-learn, Notebooks offers [What-if Tool \(WIT\)](#) and [Language Interpretability Tool \(LIT\)](#).

Best practices for using Vertex AI TensorBoard

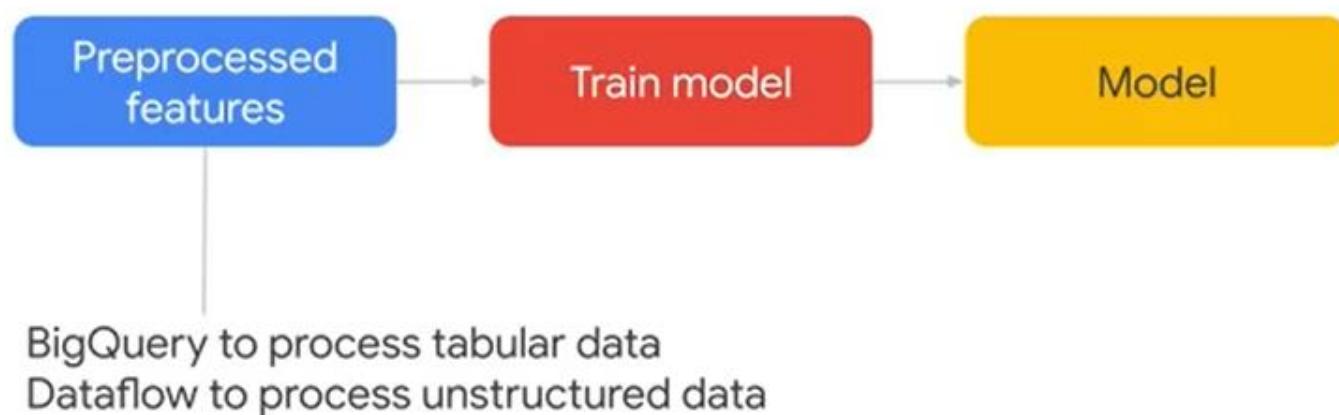
TensorBoard

Use Vertex AI TensorBoard to visualize experiments.

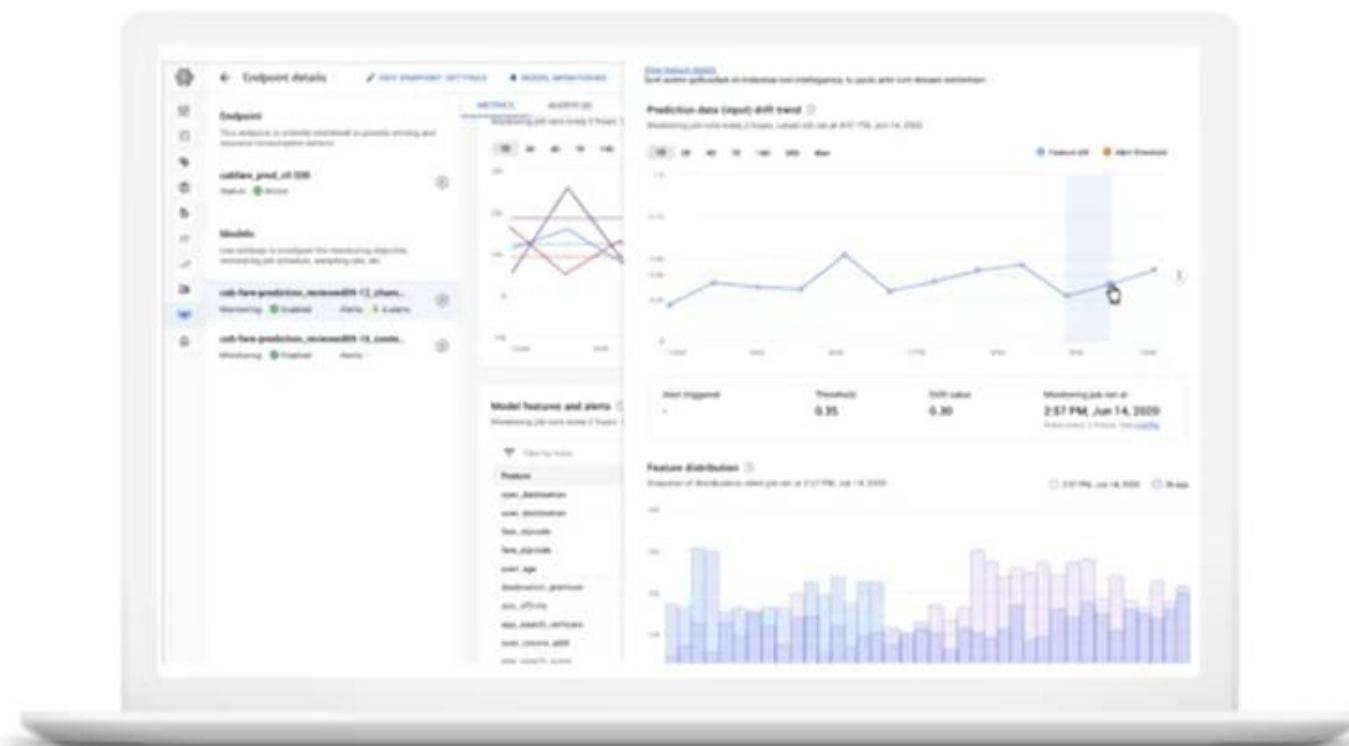


[Vertex AI TensorBoard](#) service lets you track experiment metrics such as loss and accuracy over time, visualize a model graph, project embeddings to a lower dimensional space, and much more.

For training and evaluation, we need preprocessed features



Best practices: Data preprocessing



Dataflow

Use Dataflow to process unstructured data.

TensorFlow Extended

Use TensorFlow Extended when leveraging TensorFlow ecosystem.

Data preprocessing with BigQuery

BigQuery

Use BigQuery to process tabular data.

If you're using tabular data, use BigQuery for data processing and transformation steps.

When you're working with ML, use BigQuery ML in BigQuery. Perform the transformation as a normal BigQuery query, then save the results to a [permanent table](#).

Using managed datasets in Vertex AI

Managed datasets

Use managed datasets to link data to your models.

Managed datasets:

- Enable you to create a clear link between your data and custom-trained models,
- Provide descriptive statistics and automatic or manual splitting into train, test, and validation sets.
- Are not required to use Vertex AI.

Transforming unstructured data with Dataflow

Dataflow

Use Dataflow to process unstructured data.

Use Dataflow to convert the unstructured data into binary data formats like TFRecord, which can improve performance of data ingestion during training.

If you need to perform transformations that are not expressible in Cloud SQL or are for streaming, you can use a combination of Dataflow and the [pandas](#) library.

TensorFlow Extended

TensorFlow Extended

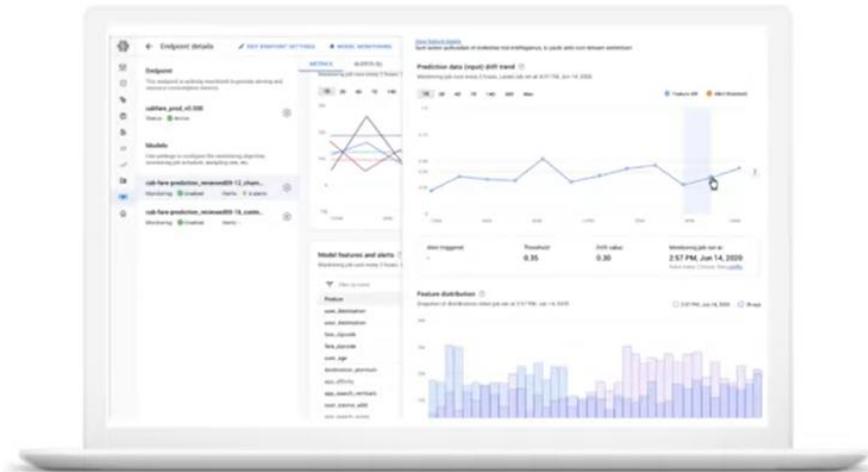
Use TensorFlow Extended when leveraging TensorFlow ecosystem.

If you're using TensorFlow for model development, use [TensorFlow Extended](#) to prepare your data for training.

[TensorFlow Transform](#) is the TensorFlow component that enables defining and executing a preprocessing function to transform your data.

Best practices for machine learning environment setup

Best practices: ML environment setup



Workbench Notebooks

Use for development and experimentation. Create NB for each team member. Use Vertex SDK for Python.

Security

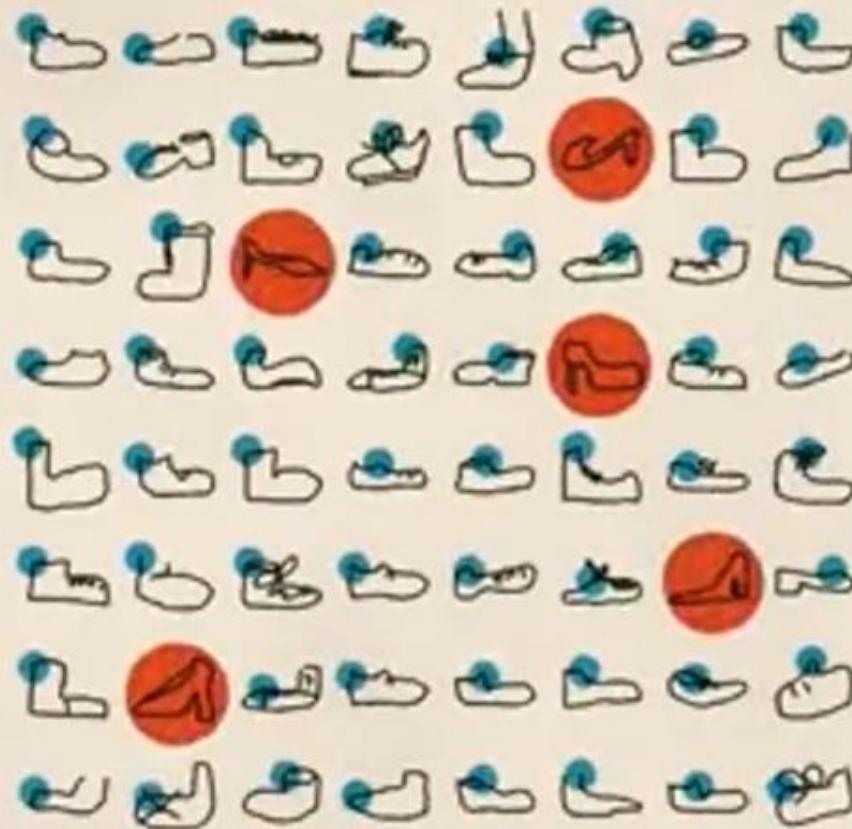
Secure PII in Notebooks.

Data & model

Store prepared data and model in same project.

Optimize performance & cost

Optimize performance and cost.



interaction bias



latent bias



selection bias

Biases in data

Unconscious biases exist in data

Examples of human biases in data

- Reporting bias
- Selection bias

Examples of human biases in collection and labeling

- Confirmation bias
- Automation bias

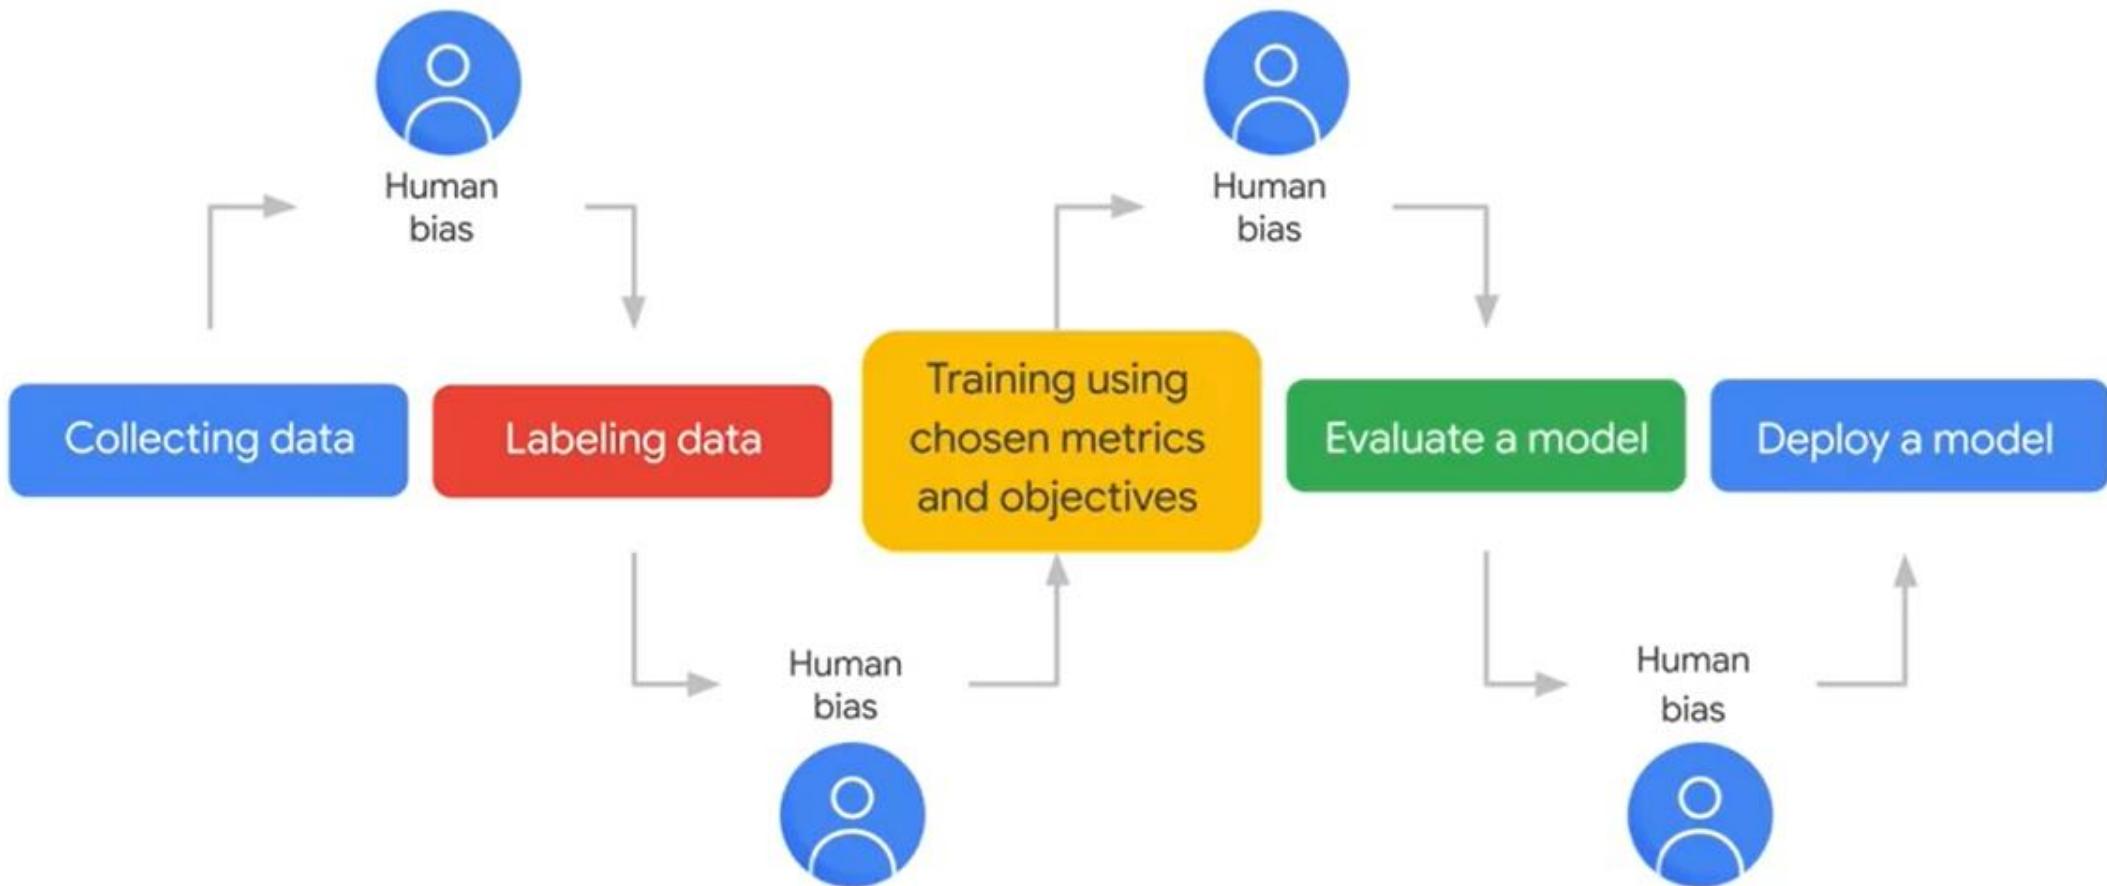
Unconscious bias from “the world” that we might reflect in ML when using existing data

Collecting data

Labeling data

Unconscious bias in our procedures that we might reflect in our ML

A typical ML pipeline with bias



Avoid creating or reinforcing unfair bias

ML models learn from existing data collected from the real world, and so an accurate model may learn or even amplify problematic pre-existing biases in the data based on race, gender, religion, or other characteristics.

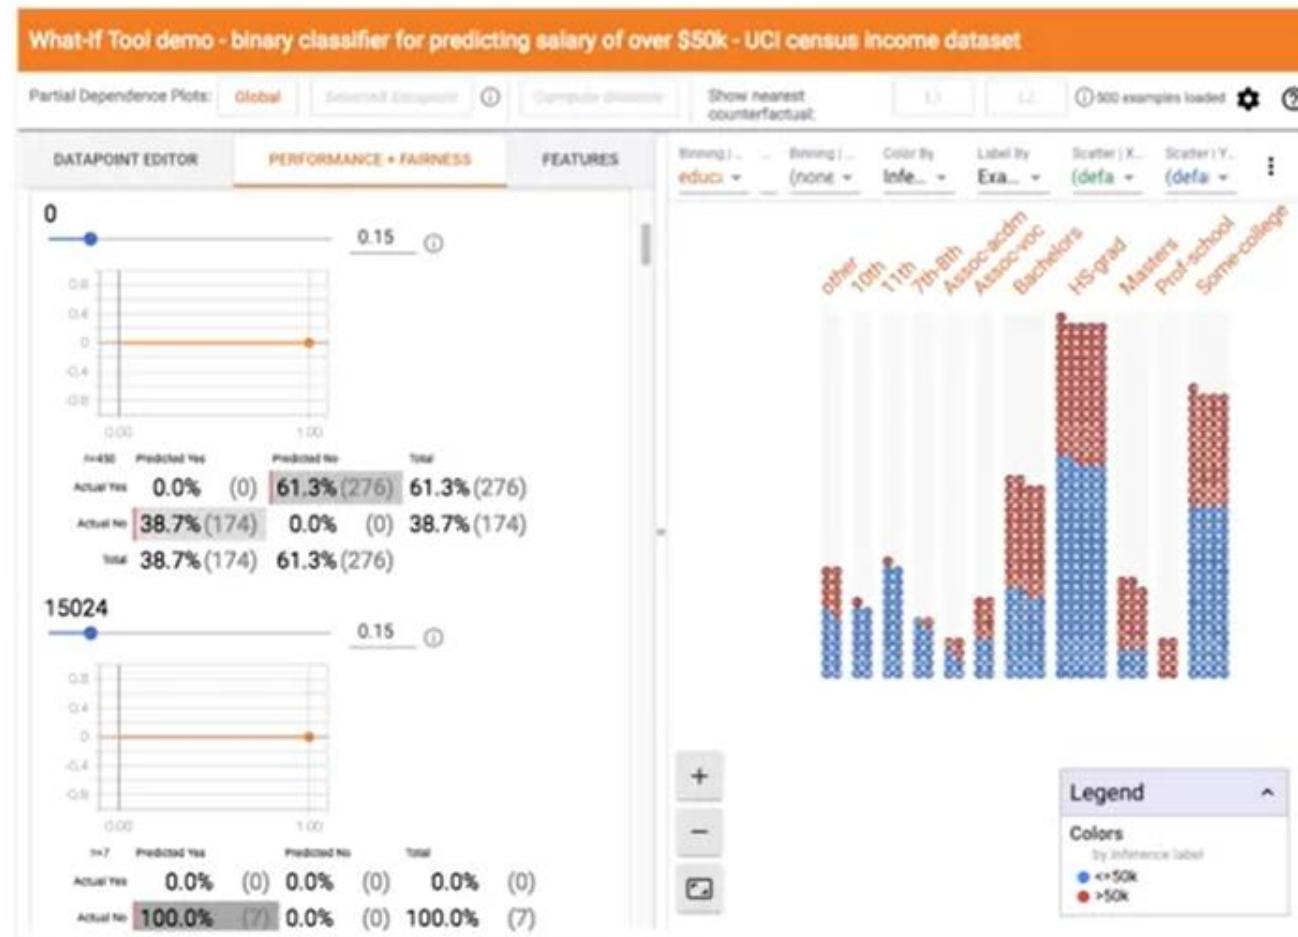
ai.google/principles



A checklist for bias-related issues

- Biometrics
- Race
- Skin color
- Religion
- Sexual orientation
- Socioeconomic status
- Income
- Country
- Location
- Health
- Language
- Dialect

Tools for responsible AI



False positives and false negatives errors occur when predictions and labels disagree

		Model predictions	
		Positive	Negative
Labels	Positive	True positives (TP)	
	Negative	False positives (FP) Type I error	
		Model says: yes	Model says: no
Labels	Positive	False negatives (FN) Type II error	
	Negative	True negatives (TN)	

How to find errors in your dataset using Facets